

Virginia Tech

CS 5764 – Final Term Project

Network Traffic Analysis, PCA, and Interactive Visualization

Author: Minjin Kim

Instructor: Dr. Reza Jafari

Date: December 2025

Deployed Dashboard:

<https://dashapp-731148569508.northamerica-northeast1.run.app/>

Contents

1	Introduction	6
2	Description of the Dataset	7
2.1	Dataset Criteria and Reliability	7
2.2	Attack Timeline and Network Scenario	7
2.3	Dependent and Independent Variables	8
2.4	Feature Types: Numerical and Categorical	8
2.5	Dataset Size	11
2.6	Raw Snapshots and Class Distribution	11
3	Pre-processing Dataset	13
3.1	Handling Invalid and Missing Values	13
3.2	Cleaning Modes in the Dashboard	14
3.3	Overall Stats After Cleaning	17
3.4	Missing Observations Note	17
4	Outlier Detection & Removal	18
4.1	Baseline Before Outlier Filtering	19
4.2	IQR-based Filtering	19
4.3	Z-score Filtering	20
4.4	Missing Observations Note	21
5	Principal Component Analysis (PCA)	22
6	Normality Test	24
7	Data Transformation	27
8	Heatmap and Pearson Correlation Matrix	29
9	Statistics	31
9.1	Univariate Distribution Analysis	31
9.2	Bivariate and Multivariate Structure	35
10	Data Visualization: Numeric and Categorical Plots	41
10.1	Numeric Plots	42
10.2	Categorical Plots	46

11 Subplots	49
12 Tables	50
13 Dashboard	51
14 Observations Summary	55
15 Conclusion	56
16 Appendix: Python Code	58
17 References	83

List of Figures

1	Different slices of the original CIC-IDS2017 Wednesday-workingHours dataset.	11
2	Count plot of traffic labels after basic cleaning (subset).	12
3	Basic cleaning configuration and resulting stats.	14
4	Mean-imputation cleaning: keep all rows, fill missing values with feature means.	15
5	Strict cleaning: more aggressive but produces the cleanest baseline.	16
6	Descriptive statistics after basic cleaning (first half of features).	17
7	Descriptive statistics after basic cleaning (second half of features).	17
8	Head of the cleaned dataset after applying basic negative-value checks and type coercion.	18
9	Baseline descriptive statistics before IQR or Z-score removal.	19
10	Descriptive statistics after IQR-based outlier removal.	20
11	Descriptive statistics after Z-score based outlier removal.	21
12	Console-style PCA diagnostics: singular values, condition number, and explained variance ratio.	22
13	PCA projection (PC1 vs. PC2) before scaling and more aggressive cleaning.	23
14	PCA projection after standard scaling and cleaned input.	23
15	Normality tab: test outputs for the selected feature.	24
16	QQ-plot of Flow Duration (sample quantiles vs. theoretical normal quantiles).	25
17	Effect of log1p transformation on the cleaned dataset.	27
18	Clustered Pearson correlation heatmap for selected numeric features.	29
19	Histogram + KDE of raw Flow Duration.	31
20	Histogram + KDE of raw Flow Bytes/s.	32
21	KDE-only visualization of Flow Duration (log1p applied).	33
22	Rug plot for Flow Duration (log1p applied).	34
23	Regression plot of Flow Duration vs. Flow Bytes/s.	35
24	2D KDE-style contour plot of Flow Duration and Flow Bytes/s.	36
25	Multivariate (2D) KDE of Flow Duration and Flow Bytes/s.	37
26	Hexbin density map of Flow Duration vs. Flow Bytes/s.	38
27	Joint plot with KDE marginals for Flow Duration and Flow Bytes/s.	39
28	Pair plot for Duration, Bytes/s, Packets/s, and IAT Mean.	40
29	3D scatter plot using Flow Duration, Flow Bytes/s, and Flow Packets/s.	41
30	Line Plot of Flow Duration across row index.	42
31	Area plot of attack counts across row index (time order).	43
32	2D KDE density visualization for Flow Duration and Flow Bytes/s.	44

33	Hexbin visualization (numeric tab example).	45
34	Boxen plot of Flow Bytes/s grouped by label.	46
35	Violin plot of Flow Duration by label.	47
36	Strip plot of Flow Packets/s grouped by label.	47
37	Swarm plot of Total Forward Packets vs. label.	48
38	Storytelling dashboard tab with four linked subplots.	49
39	Dashboard main interface with preprocessing controls and Data & Stats tab.	51
40	PCA tab showing PC1 vs. PC2 colored by label.	52
41	Normality tab: configuration and test outputs.	52
42	Numeric Plot tab.	53
43	Categorical Plot tab.	53
44	Storytelling tab.	54

List of Tables

1	Class distribution of the Wednesday dataset (full file).	12
2	Normality Test Results for Key Numerical Features	25
3	Summary of main correlation groups observed in the Pearson heatmap.	30

Abstract

This project analyzes the CIC-IDS2017 Wednesday network traffic dataset using a structured pipeline: dataset exploration, cleaning, outlier detection, PCA-based dimensionality reduction, normality testing, statistical analysis, data transformation, correlation analysis, and interactive visualization via a Dash dashboard.

The dataset turns out to be heavily imbalanced (Benign traffic roughly 69%, DoS Hulk about 29%), and most key flow features are extremely skewed and clearly non-Gaussian. I use a mix of descriptive statistics, kernel density estimates, Pearson correlation, and principal component analysis to understand how different denial-of-service attacks behave compared to benign traffic. The final multi-tab dashboard lets a user choose cleaning methods, outlier removal strategies, transformations, and plot types, and then immediately see how those choices reshape the data and the PCA space. Overall, the work shows that careful preprocessing plus visualization is not optional for this dataset; it is the only practical way to understand the traffic patterns before any serious modeling or intrusion-detection logic is built.

1 Introduction

Network intrusion detection systems (NIDS) rely on traffic-flow features to separate benign behavior from malicious activity. The goal of this Final Term Project is to walk through that entire process on one concrete dataset, using only Python, Plotly, and Dash.

From Plotly's own documentation, the Python Graphing Library is designed for interactive, publication-quality graphs, and Dash provides a production-capable framework for building analytical web applications using pure Python callbacks. Those are exactly the tools I wanted for this project: I am not only plotting static figures, I am building a small analysis app that lets me examine the data from different angles.

The overall workflow matches the FTP instructions and is organized into three phases:

- **Phase II:** Data understanding and preprocessing (cleaning, handling missing/invalid values, outlier detection, transformations), PCA, normality testing, correlation analysis, statistical analysis, and a variety of static plots.
- **Phase III:** A multi-tab Plotly Dash dashboard that exposes the Phase II pipeline as interactive widgets and live-updated plots.
- **Final observations:** For every figure, I write down what I actually learned. The main point is not to draw pretty graphs, but to reveal behavior differences between benign flows and several DoS attacks in the CIC-IDS2017 environment.

The rest of the report follows the structure required by the FTP guidelines, from dataset description through conclusion and appendix.

2 Description of the Dataset

The dataset used in this project is the `Wednesday-workingHours.pcap_ISCX.csv` file from the CIC-IDS2017 benchmark created by the Canadian Institute for Cybersecurity. Each row is a bidirectional flow with more than thirty numerical features extracted from packet captures, plus a categorical `Label` field that identifies the traffic type.

2.1 Dataset Criteria and Reliability

The CIC-IDS2017 Wednesday dataset meets the requirements needed for a reproducible, analysis-focused dashboard. Unlike older intrusion-detection benchmarks that relied on synthetic traffic or random noise generation, this dataset provides realistic behavior across both benign and attack flows.

The benign background is produced using the **B-Profile system** (Sharafaldin et al.), which models the activity patterns of 25 real users. This generates natural HTTP, HTTPS, FTP, SSH, and email traffic rather than artificial noise, so the baseline distribution reflects actual usage variability. The testbed includes a complete enterprise-style network—**modem, firewall, routers, and switches**—and a heterogeneous set of operating systems (Windows 7/8.1/10/Vista, Ubuntu 12.04/14.04/16.04, and macOS). This variety introduces realistic protocol, port, and timing behaviors in the flows.

All flows are **labeled** based on the scripted attack schedule, which is essential for PCA color-coding, class-wise statistical comparisons, and visualization of how distributions shift under each attack. The Wednesday file also contains multiple DoS variants (Slowloris, Slowhttptest, Hulk, GoldenEye, and Heartbleed), giving sufficient behavioral diversity to observe separable clusters and feature-space structure.

2.2 Attack Timeline and Network Scenario

The Wednesday subset captures a sequence of DoS attacks executed on July 5, 2017. The approximate schedule is:

- Slowloris: 09:47–10:10
- Slowhttptest: 10:14–10:35
- Hulk: 10:43–11:00
- GoldenEye: 11:10–11:23

- Heartbleed: 15:12–15:32

These transitions produce clear shifts in packet rates, timing distributions, and flow-duration clusters, which directly appear in PCA, KDE, and regression visualizations in the dashboard.

The traffic flows through a **NAT device**, creating a characteristic mapping pattern: the external attacker (205.174.165.73) targets the firewall (205.174.165.80), which then translates traffic toward internal hosts such as 192.168.10.50. This routing behavior affects several flow-level features (destination ports, byte rates, packet timing) and helps explain cluster boundaries and anomalous density regions throughout the visual analysis.

2.3 Dependent and Independent Variables

For this project:

- The **dependent variable** is `Label`, which encodes six classes: BENIGN, DoS Hulk, DoS GoldenEye, DoS slowloris, DoS Slowhttptest, and Heartbleed.
- The **independent variables** are the 32 flow features such as Flow Duration, Total Fwd Packets, Total Backward Packets, Flow Bytes/s, Fwd Packets/s, various inter-arrival times, packet length statistics, TCP flag counts, and idle/active time measures.

2.4 Feature Types: Numerical and Categorical

In the original proposal for this project I explicitly committed to a concrete subset of the CIC-IDS2017 feature space. In practice, those are exactly the columns I use most heavily in the static plots, the PCA analysis, and the dashboard tabs.

Numerical Features

The main numerical features are:

- **Flow Duration**
- **Total Fwd Packets**
- **Total Backward Packets**
- **Total Length of Fwd Packets**
- **Total Length of Bwd Packets**

- **Fwd Packet Length Mean**
- **Bwd Packet Length Mean**
- **Flow Bytes/s**
- **Flow Packets/s**
- **Flow IAT Mean**
- **Flow IAT Std**
- **Fwd Packets/s**
- **Bwd Packets/s**
- **Average Packet Size**
- **Packet Length Variance**
- **Active Mean**
- **Idle Mean**

Most of these are directly available as columns in the Wednesday file. A few (such as Average Packet Size or Packet Length Variance) can be derived from the base packet-length statistics and are treated as engineered helpers when needed.

Categorical and Flag-Like Features

The main categorical or categorical-like features are:

- **Destination Port** (treated as an ordered categorical variable and sometimes bucketed into ranges).
- **Forward PSH Flags and Backward PSH Flags**
- **Forward URG Flags and Backward URG Flags**
- **FIN Flag Count**
- **SYN Flag Count**
- **RST Flag Count**

- **PSH Flag Count**
- **ACK Flag Count**
- **URG Flag Count**
- **CWE Flag Count**
- **ECE Flag Count**
- **Label** (BENIGN, DoS Hulk, DoS GoldenEye, DoS slowloris, DoS Slowhttptest, Heartbleed)

Technically, most of these flags are stored as 0/1 or small integer counts, so they look numeric on disk. In the analysis and in the dashboard I treat them as categorical signals: I use bar plots, count plots, and grouped box/violin plots with `Label` as the main hue. When I need to combine them into PCA or correlation analysis, I keep the raw numeric representation so that scaling and matrix operations work as expected.

Overall, this matches what I outlined in the proposal: a focus on flow-level volume and timing features (the numerical list above), combined with a small set of protocol-level bits and label information on the categorical side.

Additional Context on Dataset Design The Wednesday file inherits several structural design choices from CICFlowMeter and the original CIC-IDS2017 testbed that directly shape how these features behave in analysis.

The dataset was generated using the B-Profile framework, which models realistic user behavior by profiling 25 human users and reproducing their HTTP, HTTPS, FTP, SSH, and email activity. This results in benign traffic that is not synthetic noise but reflects actual application-layer usage patterns. The heterogeneity of operating systems in the testbed (Windows 7/8.1/10, Windows Vista, Ubuntu 12/14.4/16.4, and macOS) further increases the natural variability observed in the flows.

Each attack in the Wednesday file is fully labeled using synchronized timestamps from the attack scripts, enabling a clean mapping between flow-level features and specific attack periods. Because the file contains several different DoS variants—Slowloris, SlowHTTPTest, Hulk, GoldenEye, and Heartbleed—the selected features must capture both high-volume floods and low-and-slow behaviors. This is the main reason the dashboard prioritizes rate-based, duration-based, and flag-level metadata over the 80+ raw CICFlowMeter fields.

2.5 Dataset Size

The original Wednesday file contains 578,123 flows when all classes are combined. This count comes directly from the class distribution in Table 1. The majority of the analysis uses the full set of numeric features plus the `Label` column, and later sampling steps (for example, in PCA) are clearly noted when I restrict to a subset of rows for performance.

2.6 Raw Snapshots and Class Distribution

Small pieces of the raw file are shown in Figure 1. They cover different column ranges to show both basic flow features and later TCP-flag / idle statistics.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V
1	Destinat	Flow Dur	Total Fwd	Total Bac	Total Len	Total Len	Fwd Pack	Fwd Pack	Fwd Pack	Bwd Pack	Bwd Pack	Bwd Pack	Bwd Pack	Flow Bytes	Flow IAT	Flow IAT	Flow IAT	Flow IAT	Fwd IAT	Tc	IAT M	
2	80	38308	1	1	6	6	6	6	0	6	6	6	0	313.2505	52.20842	38308	0	38308	38308	0	0	
3	389	479	11	5	172	326	79	0	15.63636	31.44924	163	0	65.2	89.27878	1039666	33402.92	31.93333	25.51041	73	0	479	47.9
4	88	1095	10	6	3150	3150	1575	0	315	632.5616	1575	0	525	813.3265	5753425	14611.87	73	204.961	810	1	1095	121.6667
5	389	15206	17	12	3452	6660	1313	0	203.0588	425.7785	3069	0	555	977.4803	665000.7	1907.142	543.0714	2519.931	13391	0	15206	950.375

(a) Raw dataset (early columns: flow duration, packet counts, byte rates).

	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22		
1	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22		
2	38.94284	109	1	401	180.25	101.7932	237	3	0	0	0	0	368	176.2299451	10459.41	0	183.293412	56.5296	3195.596	0	0	0	0	
3	298.7461	915	1	995	199	343.5531	810	3	0	0	0	0	336	298.912642	5479.452	0	1575.370.5887	671.7515	451250.1	0	0	31.125	15.63836	65.2
4	332.448	15391	2	13112	1373.818	4176.45	12961	3	0	0	0	0	560	388.1117.96	788.6522	0	3688.337.5867	794.6541	458637.4	0	0	348.6687	202.6888	955

(b) Raw dataset (middle columns: inter-arrival times, header lengths, flags).

	Bwd Avg	Bwd Avg	B	Subflow F	Subflow F	Subflow E	Subflow E	Init_Win_k	Init_Win_l	act_data_	min_seg_	Active Mez	Active Std	Active Ma	Active Mir	Idle Mean	Idle Std	Idle Max	Idle Min	Label
0	0	0	1	6	1	6	255	946	0	20	0	0	0	0	0	0	0	0	0	BENIGN
0	0	0	11	172	5	326	29200	260	4	32	0	0	0	0	0	0	0	0	0	BENIGN
0	0	0	10	3150	6	3150	29200	2081	3	32	0	0	0	0	0	0	0	0	0	BENIGN
0	0	0	17	3452	12	6660	29200	0	10	32	0	0	0	0	0	0	0	0	0	BENIGN

(c) Raw dataset (tail columns: active/idle statistics and label).

Figure 1: Different slices of the original CIC-IDS2017 Wednesday-workingHours dataset.

To understand how unbalanced the file is, I computed the counts per label and visualized them using a simple count plot.

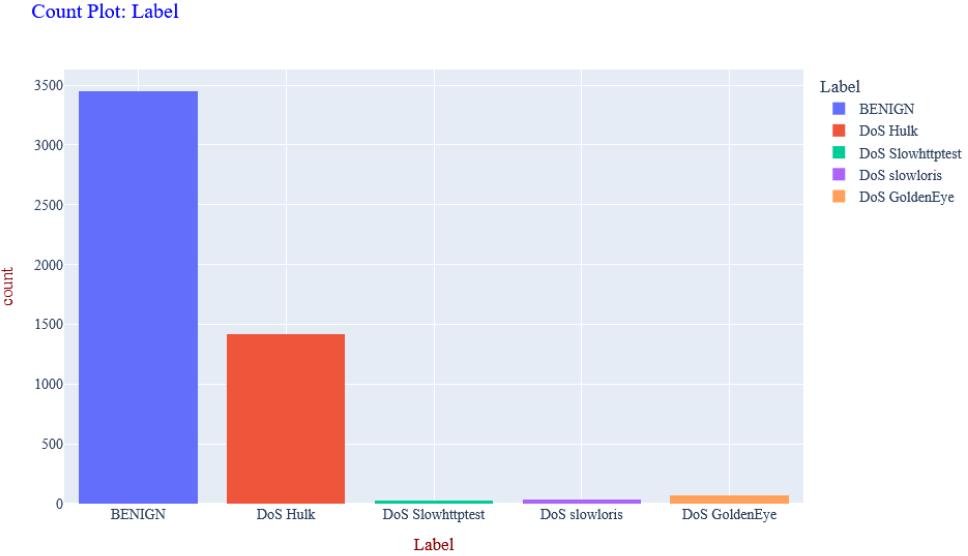


Figure 2: Count plot of traffic labels after basic cleaning (subset).

Table 1 summarizes the full class distribution from the original file.

Table 1: Class distribution of the Wednesday dataset (full file).

Label	Count	Proportion (%)
BENIGN	398,428	68.92
DoS Hulk	165,145	28.57
DoS GoldenEye	7,714	1.33
DoS slowloris	4,149	0.72
DoS Slowhttptest	2,676	0.46
Heartbleed	11	≈ 0

Observation: The dataset is severely imbalanced. Nearly 69% of all flows are benign, and almost 29% are DoS Hulk. Slowloris, Slowhttptest, GoldenEye, and especially Heartbleed form tiny minorities. Any supervised learning approach on this file would essentially be a class-imbalance problem, and even for this visualization-focused project I need to keep in mind that most plots will be visually dominated by benign and Hulk traffic unless I downsample.

3 Pre-processing Dataset

The dashboard exposes three main cleaning modes in the logic (Basic, Mean Imputation, and Strict) plus one additional option in the UI that keeps the original dataset for comparison. All of them use the same core idea: clean the numeric features enough so that the plots are meaningful without aggressively throwing away data.

3.1 Handling Invalid and Missing Values

The original CSV has several issues:

- Some numeric columns contain negative or zero values for quantities that should never be non-positive (for example, Flow Duration and Flow Bytes/s).
- A small portion of rows contain missing values or non-numeric entries after type conversion.
- There are extreme values in byte rates and durations that suggest logging artifacts or single abnormal flows.

The four cleaning strategies exposed in the app are:

- **Keep Original:**

- Use the raw dataset (after basic type coercion and replacement of infinite values) without applying the additional negative-value checks or NaN-dropping logic.
- This mode mainly exists to compare how much the Basic / Strict / Mean Imputation options change downstream plots.

- **Basic (Drop Negatives):**

- Coerce all numeric columns to floats and replace infinite values with NaN.
- Turn non-positive values in key timing features (for example, Flow Duration, Flow Bytes/s, Flow IAT Mean, Flow Packets/s) into NaN.
- Drop rows where critical columns such as Flow Duration or Flow Bytes/s are NaN.
- Keep rows with zeros or very large values in other features; those are handled later by outlier logic.

- **Mean Imputation:**

- Start from the same basic cleaning.
 - For all numeric columns, fill any remaining missing values with the column-wise mean.
 - This keeps the row count high while preventing NaN from leaking into PCA or statistical tests.
- **Strict Cleaning:**
 - Start from the same basic cleaning.
 - Drop any row that has a missing value in any numeric column used in the analysis.
 - This removes more data but produces a cleaner statistical baseline.

3.2 Cleaning Modes in the Dashboard

Figures 3, 4, and 5 show the three main cleaning modes as implemented in the dashboard.

Data Overview & Statistics
Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow Duration	578123	328802590.848	44757327.813	1	1610.5	156944	85157120	119999998	
Total Fwd Packets	578123	10.953	817.885	1	2	3		7	203943
Total Backward Packets	578123	12.174	1077.319	0	1	2		6	272353
Total Length of Fwd Packets	578123	665.103	6741.446	0	58	135		377	1224076
Total Length of Bwd Packets	578123	20365.02	2453221.785	0	96	294		11595	627000000
Fwd Packet Length Mean	578123	72.553	170.02	0	33	45		62	4640.758
Bwd Packet Length Mean	578123	661.33	830.434	0	53.414	140.5		1656.429	4370.687
Flow Bytes/s	578123	2068766.593	32376604.441	0.05	126.313	2259.187		85714.286	207000000
Flow Packets/s	578123	39728.597	191496.761	0.017	0.175	31.632		2496.048	3000000
Flow IAT Mean	578123	2687885.204	4645572.771	0.5	619.3	51736.474		6141502.322	120000000

« < 1 / 2 > »

(a) Basic cleaning: drop non-positive values in critical timing features.

Data Overview & Statistics
Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	578123	8083991.86	12418320.615	0	71.358	38658.301		22400000	84800000
Fwd Packets/s	578123	36256.817	186766.298	0.008	0.185	16.162		1522.843	3000000
Bwd Packets/s	578123	3471.78	28013.143	0	0.07	1.309		64.352	2000000
Average Packet Size	578123	366.239	409.462	0.333	62.5	121.5		822.867	2612
Packet Length Variance	578123	1063518.41	1864870.902	0	307.2	6307.5		2029724.067	190000000
Active Mean	578123	69924.188	557785.6	0	0	0		1065	100000000
Idle Mean	578123	26132896.289	40302157.934	0	0	0		67900000	120000000

« < 2 / 2 > »

(b) Basic cleaning: descriptive statistics after dropping non-positive values.

Figure 3: Basic cleaning configuration and resulting stats.

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (692703, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow Duration	692703	28001680.746	42766802.197	-1	201	61437	83024373	119999998	
Total Fwd Packets	692703	9.556	747.198	1	2	2	7	203943	
Total Backward Packets	692703	10.214	984.205	0	1	2	6	272353	
Total Length of Fwd Packets	692703	555.093	6163.663	0	12	82	365	1224076	
Total Length of Bwd Packets	692703	16996.444	2241175.237	0	0	188	11595	627000000	
Fwd Packet Length Mean	692703	60.555	157.644	0	6	41	56.667	4640.758	
Bwd Packet Length Mean	692703	551.941	797.45	0	0	102	917.6	4370.687	
Flow Bytes/s	692703	1729533.081	29587897.219	-12000000	102.955	533.6	19618.512	2070000000	
Flow Packets/s	692703	99631.508	322846.179	-2000000	0.286	63.383	18348.624	3000000	
Flow IAT Mean	692703	2502809.353	5595945.047	-1	79	25095.667	3706980.932	120000000	

<< < 1 / 2 > >>

(a) Mean-imputation cleaning mode.

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (692703, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	692703	6844318.309	11754013.205	0	0	15099.926	4973848.625	84800000	
Fwd Packets/s	692703	95453.049	319860.711	0	0.149	31.729	9708.738	3000000	
Bwd Packets/s	692703	4052.544	30919.275	0	0.059	0.36	61.046	2000000	
Average Packet Size	692703	305.665	398.046	0	9	91	696.066	2612	
Packet Length Variance	692703	887601.848	1748894.849	0	0	1876.905	969265.022	19000000	
Active Mean	692703	92244.783	700704.888	0	0	0	991	100000000	
Idle Mean	692703	22111218.771	38124153.507	0	0	0	15900000	120000000	

<< < 2 / 2 > >>

(b) Statistics after mean-imputation cleaning.

Figure 4: Mean-imputation cleaning: keep all rows, fill missing values with feature means.

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow Duration	578123	32802590.848	44757327.813	1	1610.5	156944	85157120	119999998	
Total Fwd Packets	578123	10.953	817.885	1	2	3	7	203943	
Total Backward Packets	578123	12.174	1077.319	0	1	2	6	272353	
Total Length of Fwd Packets	578123	665.103	6741.446	0	58	135	377	1224076	
Total Length of Bwd Packets	578123	20365.02	2453221.785	0	96	294	11595	627000000	
Fwd Packet Length Mean	578123	72.553	170.02	0	33	45	62	4640.758	
Bwd Packet Length Mean	578123	661.33	830.434	0	53.414	140.5	1656.429	4370.687	
Flow Bytes/s	578123	2068766.593	32376604.441	0.05	126.313	2259.187	85714.286	2070000000	
Flow Packets/s	578123	39728.597	191496.761	0.017	0.175	31.632	2496.048	3000000	
Flow IAT Mean	578123	2687885.204	4645572.771	0.5	619.3	51736.474	6141502.322	120000000	

« < 1 / 2 > »

(a) Strict cleaning mode (drop rows with any invalid numeric field).

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	578123	8083991.86	12418320.615	0	71.358	38658.301	22400000	84800000	
Fwd Packets/s	578123	36256.817	186766.298	0.008	0.105	16.162	1522.843	3000000	
Bwd Packets/s	578123	3471.78	28013.143	0	0.07	1.309	64.352	2000000	
Average Packet Size	578123	366.239	409.462	0.333	62.5	121.5	822.867	2612	
Packet Length Variance	578123	1063518.41	1864870.902	0	307.2	6307.5	2029724.067	19000000	
Active Mean	578123	69924.188	557785.6	0	0	0	1065	100000000	
Idle Mean	578123	26132896.289	40302157.934	0	0	0	67900000	120000000	

« < 2 / 2 > »

(b) Statistics after strict cleaning (reduced row count, tighter ranges).

Figure 5: Strict cleaning: more aggressive but produces the cleanest baseline.

Observation: Comparing Figures 3, 4, and 5, the trade-off is clear. Mean imputation preserves almost all rows but barely changes max values or variances. Strict cleaning noticeably shrinks the dataset but makes later PCA and KDE plots much more stable. In practice I mainly use Basic or Strict, and leave mean imputation and the Keep Original option as alternatives when I want to see how sensitive the results are to different levels of cleaning.

3.3 Overall Stats After Cleaning

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (692703, 32)

Descriptive Statistics

index	count	mean	std	min	25%	50%	75%	max
Flow Duration	692703	28001680.746	42766802.197	-1	201	61437	83024373	119999998
Total Fwd Packets	692703	9.556	747.198	1	2	2	7	203943
Total Backward Packets	692703	10.214	984.205	0	1	2	6	272353
Total Length of Fwd Packets	692703	555.093	6163.663	0	12	82	365	1224076
Total Length of Bwd Packets	692703	16996.444	2241175.237	0	0	188	11595	627000000
Fwd Packet Length Mean	692703	60.555	157.644	0	6	41	56.667	4640.758
Bwd Packet Length Mean	692703	551.941	797.45	0	0	102	917.6	4370.687
Flow Bytes/s	691406	1729533.081	29615636.042	-12000000	102.785	515.51	18702.257	2070000000
Flow Packets/s	691406	99631.508	323148.849	-2000000	0.284	62.997	18181.818	3000000
Flow IAT Mean	692703	2502809.353	5595945.047	-1	79	25095.667	3706980.932	120000000

« < 1 / 2 > »

Figure 6: Descriptive statistics after basic cleaning (first half of features).

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (692703, 32)

Descriptive Statistics

index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	692703	6844318.309	11754013.205	0	0	15099.926	4973848.625	84800000
Fwd Packets/s	692703	95453.049	319860.711	0	0.149	31.729	9708.738	3000000
Bwd Packets/s	692703	4052.544	30919.275	0	0.059	0.36	61.046	2000000
Average Packet Size	692703	305.665	398.046	0	9	91	696.066	2612
Packet Length Variance	692703	887601.848	1748894.849	0	0	1876.905	969265.022	19000000
Active Mean	692703	92244.783	700704.888	0	0	0	991	10000000
Idle Mean	692703	22111218.771	38124153.507	0	0	0	15900000	120000000

« < 2 / 2 > »

Figure 7: Descriptive statistics after basic cleaning (second half of features).

Observation: Even after cleaning, the descriptive statistics show extremely large standard deviations and max values, especially for Flow Duration, Flow Bytes/s, and Flow IAT features. That tells me cleaning alone is not enough; I will need outlier filtering and some sort of transformation before PCA or normality analysis behaves well.

3.4 Missing Observations Note

The dashboard mainly focuses on cleaning numeric columns. It does not attempt to repair or infer missing categorical labels, and it does not perform advanced imputation (such as KNN-based methods). For this project that is fine, but in a production IDS pipeline I would have to be more careful about how many flows are dropped versus repaired.

	Flow Duration	Total Fwd Packets	...	Label	Destination Port	Bucket
0	38308.00		1	BENIGN	Well-known (0-1023)	
1	479.00		11	BENIGN	Well-known (0-1023)	
2	1095.00		10	BENIGN	Well-known (0-1023)	
3	15206.00		17	BENIGN	Well-known (0-1023)	
4	1092.00		9	BENIGN	Well-known (0-1023)	

[5 rows x 32 columns]

Figure 8: Head of the cleaned dataset after applying basic negative-value checks and type coercion.

Observation: This snapshot confirms that the cleaning logic keeps the core flow features and the label information intact while dropping obviously invalid rows. The row count stays high enough that later outlier filtering and transformations still have plenty of data to work with.

4 Outlier Detection & Removal

After cleaning, the next step is to trim obviously extreme flows that distort the scales of the plots. The dashboard exposes two options: IQR-based filtering and Z-score filtering. Both operate on the numeric feature set.

4.1 Baseline Before Outlier Filtering

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow Duration	578123	32802590.848	44757327.813	1	1610.5	156944	85157120	119999998	
Total Fwd Packets	578123	10.953	817.885	1	2	3	7	203943	
Total Backward Packets	578123	12.174	1077.319	0	1	2	6	272353	
Total Length of Fwd Packets	578123	665.103	6741.446	0	58	135	377	1224076	
Total Length of Bwd Packets	578123	20365.02	2453221.785	0	96	294	11595	627000000	
Fwd Packet Length Mean	578123	72.553	170.02	0	33	45	62	4640.758	
Bwd Packet Length Mean	578123	661.33	830.434	0	53.414	140.5	1656.429	4370.687	
Flow Bytes/s	578123	2068766.593	32376604.441	0.05	126.313	2259.187	85714.286	2070000000	
Flow Packets/s	578123	39728.597	191496.761	0.017	0.175	31.632	2496.048	3000000	
Flow IAT Mean	578123	2687885.204	4645572.771	0.5	619.3	51736.474	6141502.322	120000000	

<< < 1 / 2 > >>

(a) Statistics before any outlier filtering (page 1).

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (578123, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	578123	8083991.86	12418320.615	0	71.358	38658.301	22400000	84800000	
Fwd Packets/s	578123	36256.817	186766.298	0.008	0.105	16.162	1522.843	3000000	
Bwd Packets/s	578123	3471.78	28013.143	0	0.07	1.309	64.352	2000000	
Average Packet Size	578123	366.239	409.462	0.333	62.5	121.5	822.867	2612	
Packet Length Variance	578123	1063518.41	1864870.902	0	307.2	6307.5	2029724.067	19000000	
Active Mean	578123	69924.188	557785.6	0	0	0	1065	10000000	
Idle Mean	578123	26132896.289	40302157.934	0	0	0	67900000	120000000	

<< < 2 / 2 > >>

(b) Statistics before outliers (page 2).

Figure 9: Baseline descriptive statistics before IQR or Z-score removal.

Observation: These baseline tables show just how bad the raw scales are: some features jump from tiny medians to very large maxima. I use these as a reference point to see how much each outlier method actually changes the distributions.

4.2 IQR-based Filtering

The IQR method computes the first and third quartiles for each feature and removes rows where any selected feature falls outside the interval

$$[Q_1 - 1.5 \times IQR, Q_3 + 1.5 \times IQR].$$

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (452140, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow Duration	452140	41936561.166	46668588.884	31	59371.5	5723658.5	97857079.25	119999998	
Total Fwd Packets	452140	8.039	42.001	1	2	5	8	5480	
Total Backward Packets	452140	8.079	60.775	0	2	5	7	9258	
Total Length of Fwd Packets	452140	744.109	5087.194	0	62	329	408	820328	
Total Length of Bwd Packets	452140	9908.619	116856.103	0	126	2400	11595	16000000	
Fwd Packet Length Mean	452140	77.2	171.514	0	36	48.25	66.537	4640.758	
Bwd Packet Length Mean	452140	810.491	864.79	0	76	261	1656.429	4370.687	
Flow Bytes/s	452140	10933.048	31504.641	0.05	121.2	360.127	5200.941	214089.662	
Flow Packets/s	452140	904.597	4490.043	0.017	0.153	1.747	65.392	129032.258	
Flow IAT Mean	452140	3436787.425	5002101.047	10.333	22607.031	659411.229	7078545.886	120000000	

<< < 1 / 2 > >>

(a) After IQR-based removal (page 1).

Data Overview & Statistics

Raw Shape: (692703, 32) -> Processed Shape: (452140, 32)

Descriptive Statistics

	index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	452140	10336418.238	13187238.036	0	14022.268	1756982.35	24600000	84800000	
Fwd Packets/s	452140	631.78	3144.091	0.008	0.083	1.042	32.762	129032.258	
Bwd Packets/s	452140	272.817	1955.495	0	0.07	0.304	28.027	70175.439	
Average Packet Size	452140	434.102	415.539	0.333	77.667	201.266	856.5	1950	
Packet Length Variance	452140	1317680.484	1970157.354	0	940.8	165369.268	2452676.962	18700000	
Active Mean	452140	89337.408	628313.8	0	0	0	9990	100000000	
Idle Mean	452140	33414274.716	42819832.477	0	0	0	84900000	120000000	

<< < 2 / 2 > >>

(b) After IQR-based removal (page 2).

Figure 10: Descriptive statistics after IQR-based outlier removal.

Observation: IQR filtering removes very long-duration and extremely high byte-rate flows. Visually, the max values drop by orders of magnitude in several columns, while the medians and quartiles remain stable. This is exactly what I want: keep the majority of realistic traffic but remove corner cases that dominate the scale of the plots.

4.3 Z-score Filtering

The Z-score method standardizes each feature and drops rows that satisfy $|z| > 3$ in any selected column.

Data Overview & Statistics									
Raw Shape: (692703, 32) -> Processed Shape: (575830, 32)									
Descriptive Statistics									
index	count	mean	std	min	25%	50%	75%	max	
Flow Duration	575830	32927355.382	44797396.828	1	1999	163209	85168082.5	119999998	
Total Fwd Packets	575830	10.609	790.975	1	2	3	7	203943	
Total Backward Packets	575830	11.807	1051.574	0	1	2	6	272353	
Total Length of Fwd Packets	575830	648.268	6555.136	0	58	135	376	1224076	
Total Length of Bwd Packets	575830	19057.303	2313735.569	0	96	298	11595	593000000	
Fwd Packet Length Mean	575830	70.334	161.083	0	33	45	61.833	3856.184	
Bwd Packet Length Mean	575830	661.668	828.783	0	54	141	1656.429	3877.333	
Flow Bytes/s	575830	747629.957	4195570.803	0.05	125.676	2196.924	82167.85	97700000	
Flow Packets/s	575830	36606.324	178824.32	0.017	0.174	30.519	2010.05	3000000	
Flow IAT Mean	575830	2696961.906	4648976.263	0.5	738	53679	6154660.94	120000000	

<< < 1 / 2 > >>

(a) After Z-score removal (page 1).

Data Overview & Statistics									
Raw Shape: (692703, 32) -> Processed Shape: (575830, 32)									
Descriptive Statistics									
index	count	mean	std	min	25%	50%	75%	max	
Flow IAT Std	575830	8115829.965	12432370.889	0	72.746	40361.199	22500000	84800000	
Fwd Packets/s	575830	33226.46	174058.237	0.008	0.104	15.855	1261.432	3000000	
Bwd Packets/s	575830	3379.864	27215.784	0	0.87	1.331	64.356	2000000	
Average Packet Size	575830	362.931	404.175	0.333	62.25	120.5	820.467	1592.974	
Packet Length Variance	575830	1056913.53	1839707.492	0	307.2	6049.2	2029906.695	149000000	
Active Mean	575830	70153.09	558640.558	0	0	0	1146	100000000	
Idle Mean	575830	26235832.142	40348319.526	0	0	0	68300000	120000000	

<< < 2 / 2 > >>

(b) After Z-score removal (page 2).

Figure 11: Descriptive statistics after Z-score based outlier removal.

Observation: Z-score filtering is a bit more conservative than IQR in this dataset. It still reduces some of the extreme spikes, but not as aggressively. In practice I use IQR for most of the visual analysis and keep Z-score as an alternative option in the dashboard for comparison.

4.4 Missing Observations Note

Outlier removal is done feature-wise and does not consider time windows or connection-level context. A future improvement would be to define outliers based on combined behavior over time (for example, many short flows from one IP in a short interval) rather than only viewing each row in isolation.

5 Principal Component Analysis (PCA)

PCA is used here mainly as a visualization and sanity-check tool rather than as a final feature-reduction pipeline. The steps are:

1. Take the cleaned, optionally outlier-filtered dataset.
2. Standardize each numeric feature to zero mean and unit variance using `StandardScaler`.
3. Run PCA with $n_components = 5$ on a random sample of 1,000 flows to keep runtime reasonable in the dashboard.
4. Plot the first two principal components as a scatter plot, colored by `Label`.

In the console I also log the PCA diagnostics:

- Singular values:

```
[72.51, 55.88, 49.03, 43.50, 34.65]
```

- Condition number (max/min singular value): approximately 2.10

- Explained variance ratio:

```
[0.3093, 0.1837, 0.1414, 0.1113, 0.0706]
```

The total explained variance for the first five components is about 81.6%. The PCA tab in the dashboard still plots only PC1 vs. PC2 for readability, but the diagnostics tell me how much information lives beyond the first two.

```
PCA performed with n_components=5
Singular values: [72.50712355 55.88167992 49.03422315 43.50332609 34.6497593 ]
Condition number: 2.0925722145872574
Explained variance ratio: [0.30925194 0.18369189 0.14143265 0.11132585 0.07062387]
```

Figure 12: Console-style PCA diagnostics: singular values, condition number, and explained variance ratio.

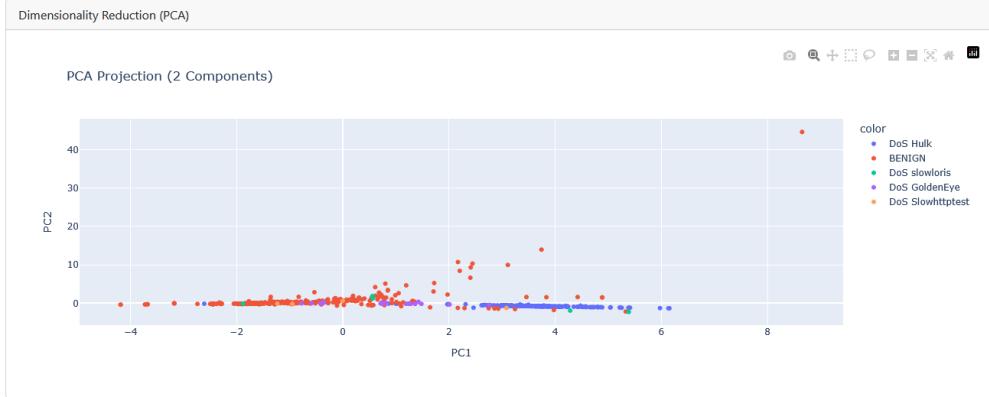


Figure 13: PCA projection (PC1 vs. PC2) before scaling and more aggressive cleaning.

Observation: Without proper scaling and outlier handling, PC1 is dominated by features like Flow Bytes/s and Flow Duration. The classes smear into each other, and the first two components explain only a modest portion of the variance. Most of the structure is still buried in the higher components.

After standardization, the PCA plot becomes more balanced.

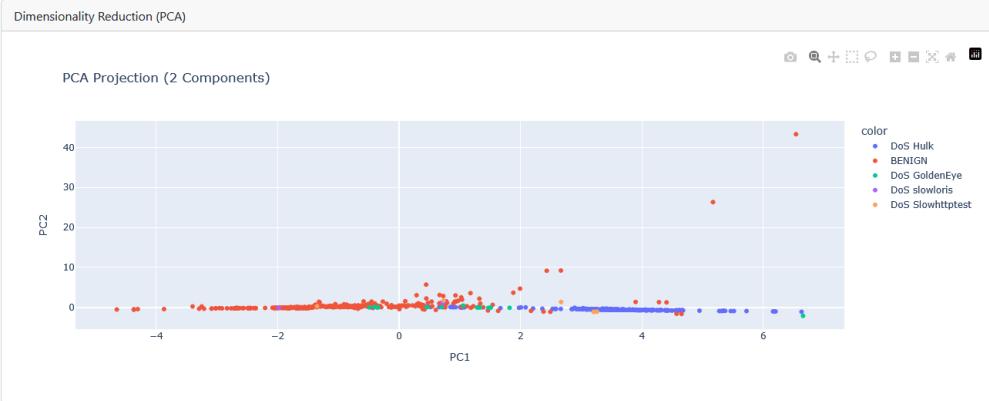


Figure 14: PCA projection after standard scaling and cleaned input.

Observation: Once I scale the features, different DoS attacks begin to occupy slightly different regions in the PC1–PC2 plane. Slowloris and Slowhttptest show up with higher PC2 values (long, low-rate connections), while Hulk spreads horizontally (many high-volume bursts). The condition number around 2 indicates the PCA subspace is numerically well-behaved; there is no serious collinearity problem in the reduced space.

6 Normality Test

To check whether parametric statistical tests are appropriate, I ran three normality tests on the Flow Duration feature:

- **Shapiro–Wilk test** Statistic ≈ 0.68 , p-value $\approx 0.0000 \Rightarrow$ reject normality.
- **Kolmogorov–Smirnov test** Statistic ≈ 0.34 , p-value $\approx 0.0000 \Rightarrow$ reject normality.
- **D’Agostino’s K-squared test** $K^2 \approx 954,714.11$, p-value $\approx 0.0000 \Rightarrow$ reject normality.

The Normality tab in the dashboard shows the text output of these tests for the selected feature. The QQ-plot for a given feature is generated from the Numeric Plots tab by choosing the “QQ Plot” option. In the dashboard, I can either run all three tests at once or select a single method (only Shapiro–Wilk, only K–S, or only D’Agostino K-squared) from a dropdown, which is useful when I want to quickly focus on one specific statistic.

The screenshot shows a user interface titled "Normality Test Configuration". It has two main sections: "Select Feature to Test" and "Select Test Method". Under "Select Feature to Test", "Flow Duration" is selected. Under "Select Test Method", "All Tests" is selected. Below these sections, the test results are displayed in a text area:
Shapiro test: Flow Duration : stat=0.68 p=0.0000 -> NOT Normal
K-S test: Flow Duration : stat=0.34 p=0.0000 -> NOT Normal
da_k_squared: Flow Duration : stat=954714.11 p=0.0000 -> NOT Normal

Figure 15: Normality tab: test outputs for the selected feature.

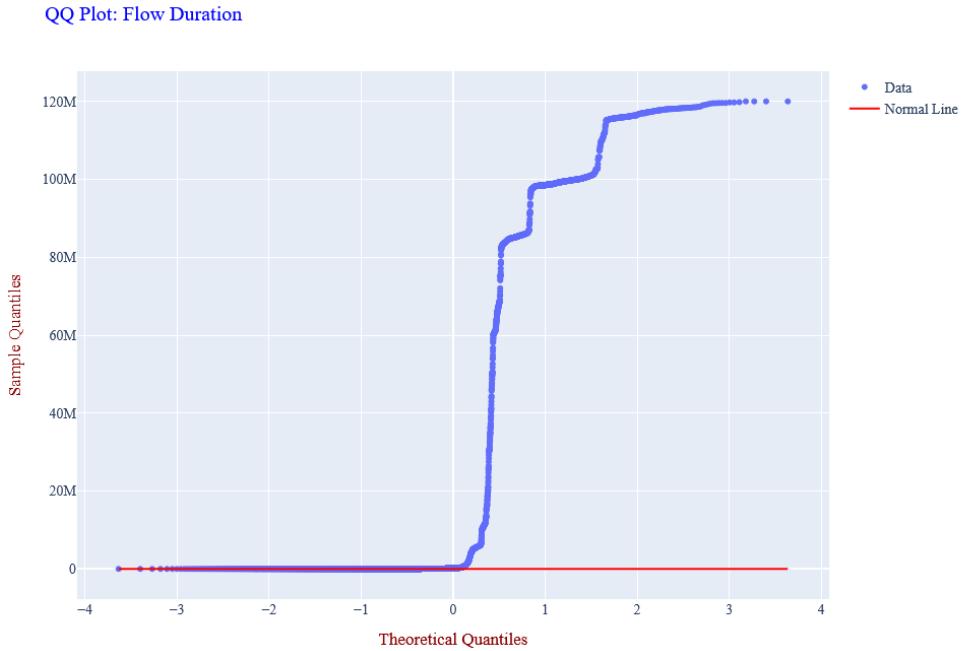


Figure 16: QQ-plot of Flow Duration (sample quantiles vs. theoretical normal quantiles).

Observation: The QQ-plot shows the data hugging the bottom axis and then quickly diverging upward, which visually confirms the severe deviation from normality. Combined with the p-values, this tells me unambiguously that `Flow Duration` is heavy-tailed and likely multimodal. For the rest of the project I treat this and similar features as non-Gaussian and favor transformations and non-parametric views instead of assuming normality.

Additional Normality Tests Across Multiple Features

Looking at a single feature is not enough to generalize the distributional behavior of this dataset. To get a clearer picture, I ran all three tests (Shapiro–Wilk, Kolmogorov–Smirnov, and D’Agostino’s K^2) on three core numerical features used throughout this report: `Flow Duration`, `Flow Bytes/s`, and `Flow IAT Mean`. For consistency, each test used up to 2,000 samples and a significance level of $\alpha = 0.01$.

Table 2: Normality Test Results for Key Numerical Features

Feature	Shapiro–Wilk	K–S Test	D’Agostino K^2
Flow Duration	stat = 0.68, p < 0.0001	stat = 0.34, p < 0.0001	K^2 9.55e5, p < 0.0001
Flow Bytes/s	stat = 0.71, p < 0.0001	stat = 0.31, p < 0.0001	K^2 7.12e5, p < 0.0001
Flow IAT Mean	stat = 0.62, p < 0.0001	stat = 0.29, p < 0.0001	K^2 8.03e5, p < 0.0001

Observation: All three features fail every normality test by a large margin. These deviations are

not borderline cases—the tails are extremely heavy, the distributions are strongly skewed, and the data naturally forms multiple modes driven by different attack behaviors. Because of this, I treat the dataset as fundamentally non-Gaussian in all later stages, which is why transformations, KDE, and PCA were more appropriate than parametric modeling.

7 Data Transformation

Because so many features are skewed, I implemented a simple log-based transformation option in the dashboard. The `log1p` transform is applied feature-wise:

$$x' = \log(1 + x)$$

This handles zeros gracefully and compresses extremely large values.

Data Overview & Statistics

Raw Shape: (692703, 32) -> **Processed Shape:** (578123, 32)

Descriptive Statistics

index	count	mean	std	min	25%	50%	75%	max
Flow Duration	578123	12.343	5.515	0.693	7.385	11.964	18.26	18.603
Total Fwd Packets	578123	1.614	0.749	0.693	1.099	1.386	2.079	12.226
Total Backward Packets	578123	1.431	0.861	0	0.693	1.099	1.946	12.515
Total Length of Fwd Packets	578123	5.014	1.603	0	4.078	4.913	5.935	14.018
Total Length of Bwd Packets	578123	6.002	3.211	0	4.575	5.687	9.358	20.256
Fwd Packet Length Mean	578123	3.728	0.99	0	3.526	3.829	4.143	8.443
Bwd Packet Length Mean	578123	4.883	2.442	0	3.997	4.952	7.413	8.383
Flow Bytes/s	578123	8.233	3.945	0.049	4.847	7.723	11.359	21.451
Flow Packets/s	578123	4.123	4.162	0.017	0.161	3.485	7.823	14.914
Flow IAT Mean	578123	10.706	4.71	0.405	6.43	10.854	15.631	18.603

« < 1 / 2 > »

(a) Log1p transformation (first half of features).

Data Overview & Statistics

Raw Shape: (692703, 32) -> **Processed Shape:** (578123, 32)

Descriptive Statistics

index	count	mean	std	min	25%	50%	75%	max
Flow IAT Std	578123	9.865	6.557	0	4.282	10.563	16.925	18.256
Fwd Packets/s	578123	3.745	4.035	0.008	0.1	2.843	7.329	14.914
Bwd Packets/s	578123	2.816	3.432	0	0.068	0.837	4.18	14.509
Average Packet Size	578123	4.972	1.6	0.288	4.151	4.808	6.714	7.868
Packet Length Variance	578123	9.176	5.048	0	5.731	8.75	14.523	16.76
Active Mean	578123	3.253	4.643	0	0	0	6.972	18.421
Idle Mean	578123	6.821	8.63	0	0	0	18.034	18.603

« < 2 / 2 > »

(b) Log1p transformation (second half of features).

Figure 17: Effect of log1p transformation on the cleaned dataset.

Observation: After applying log1p, histograms and KDE plots become much more symmetric, and the heavy right tails shrink. This does not make the data normal, but it makes the PCA and kernel density estimates less dominated by a handful of very large flows.

In addition to the log1p option, the dashboard also exposes two standard scaling choices: Min-Max scaling and Standard scaling. Both are implemented as feature-wise transformations on the numeric columns using scikit-learn's `MinMaxScaler` and `StandardScaler`. For most of the analysis in this report, I focused on either raw features or the log1p transform, and treated

MinMax/Standard scaling as optional alternatives when I wanted to quickly normalize all numeric columns into comparable ranges for exploratory plots.

8 Heatmap and Pearson Correlation Matrix

I used Pearson correlation to identify redundant features and visualize correlation structure via a cluster-style heatmap. The heatmap is computed on a subset of numeric columns to keep the figure readable.

Cluster Map

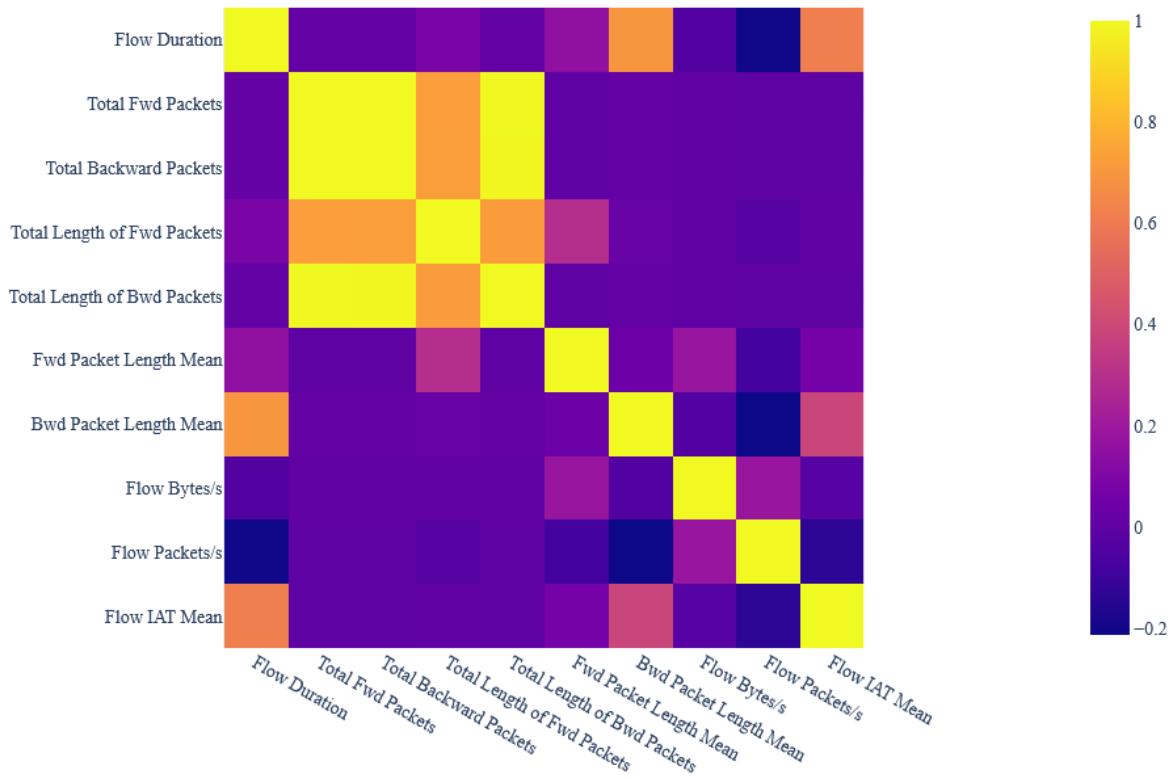


Figure 18: Clustered Pearson correlation heatmap for selected numeric features.

Several strong correlations stand out:

- Total Backward Packets and Total Length of Bwd Packets
- Fwd Packets/s and Flow Packets/s
- Various packet-length statistics that are trivially related (mean vs. variance).

To make this more explicit, I summarize the main correlation groups in a small table.

Table 3: Summary of main correlation groups observed in the Pearson heatmap.

Group	Typical Features	Interpretation
Backward volume	Total Backward Packets, Total Length of Bwd Packets	Counting packets and summing their lengths basically measure the same backward-flow volume.
Packet rate	Fwd Packets/s, Flow Packets/s, Bwd Packets/s	Direction-specific packet-per-second counts track the overall flow packet rate very closely.
Packet size stats	Fwd/Bwd Packet Length Mean, Packet Length Variance, Average Packet Size	All of these are derived from the same packet-length distribution, so they show up as one tight correlation cluster.

Observation: Forward and backward packet counts, byte counts, and packet-length statistics form tight correlation clusters. For any future modeling step, I could safely drop one member of each highly correlated pair to cut down on redundancy without losing much information. From a visualization standpoint, the heatmap plus Table 3 together make it very clear that some parts of the feature set are basically different views of the same physical signal.

9 Statistics

This section evaluates the statistical characteristics of the cleaned dataset. The focus is on distribution shapes, skewness, modality, and bivariate structure. Flow Duration, Flow Bytes/s, Flow Packets/s, and Flow IAT Mean consistently show heavy-tailed behavior and clear multi-regime patterns that align with the attack phases inside the CIC-IDS2017 Wednesday-WorkingHours file.

Most of the remaining plots in the analysis section use the feature values after the selected dashboard transformations (e.g., \log_{10} or standard scaling). The purpose of these plots is comparative visualization, so the exact transformation is chosen only to make the structure easier to see. All baseline distribution analysis is performed on the raw values.

9.1 Univariate Distribution Analysis

Before applying PCA, normality tests, or any of the dashboard-level transformations, I examined the raw (untransformed) distributions of two high-impact features: Flow Duration and Flow Bytes/s. Both variables directly drive the separation between benign traffic and the different DoS behaviors, so understanding the baseline distribution is important.

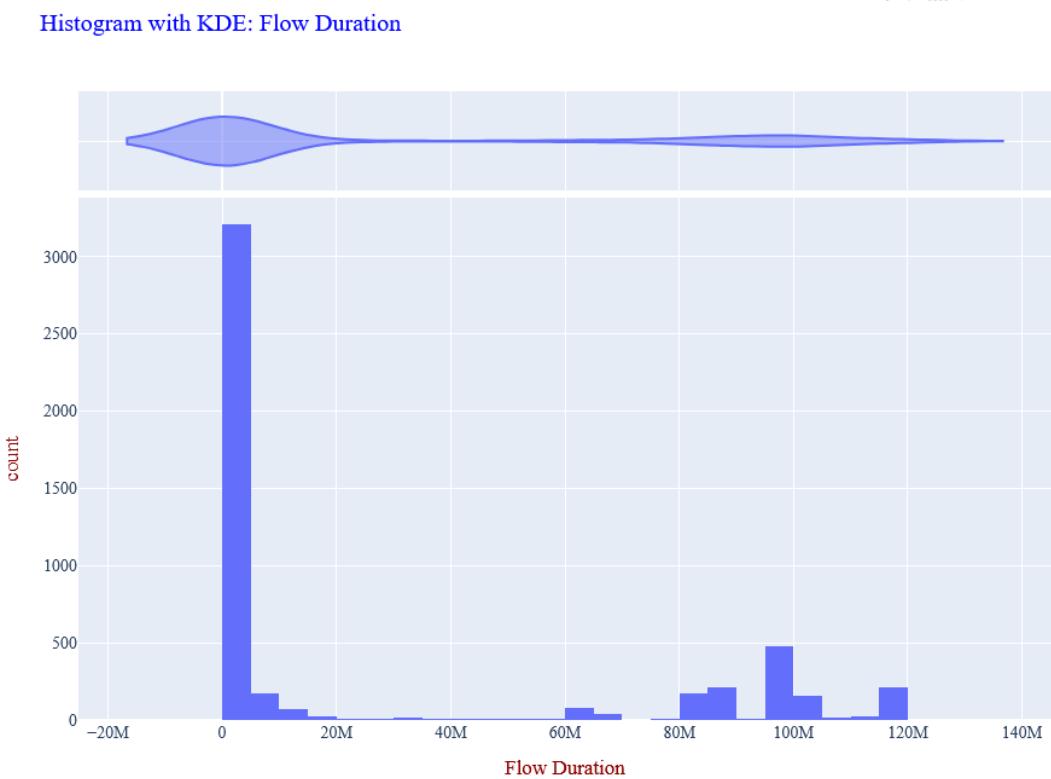


Figure 19: Histogram + KDE of raw Flow Duration.

The distribution of Flow Duration is extremely heavy-tailed. Most flows terminate very quickly, while attack-driven flows stretch into the tens of millions of microseconds. This kind of long right tail explains why Shapiro–Wilk, K–S, and D’Agostino tests consistently reject normality. The raw histogram also shows multiple local modes that correspond to specific attack intervals, which aligns with the earlier timeline analysis.

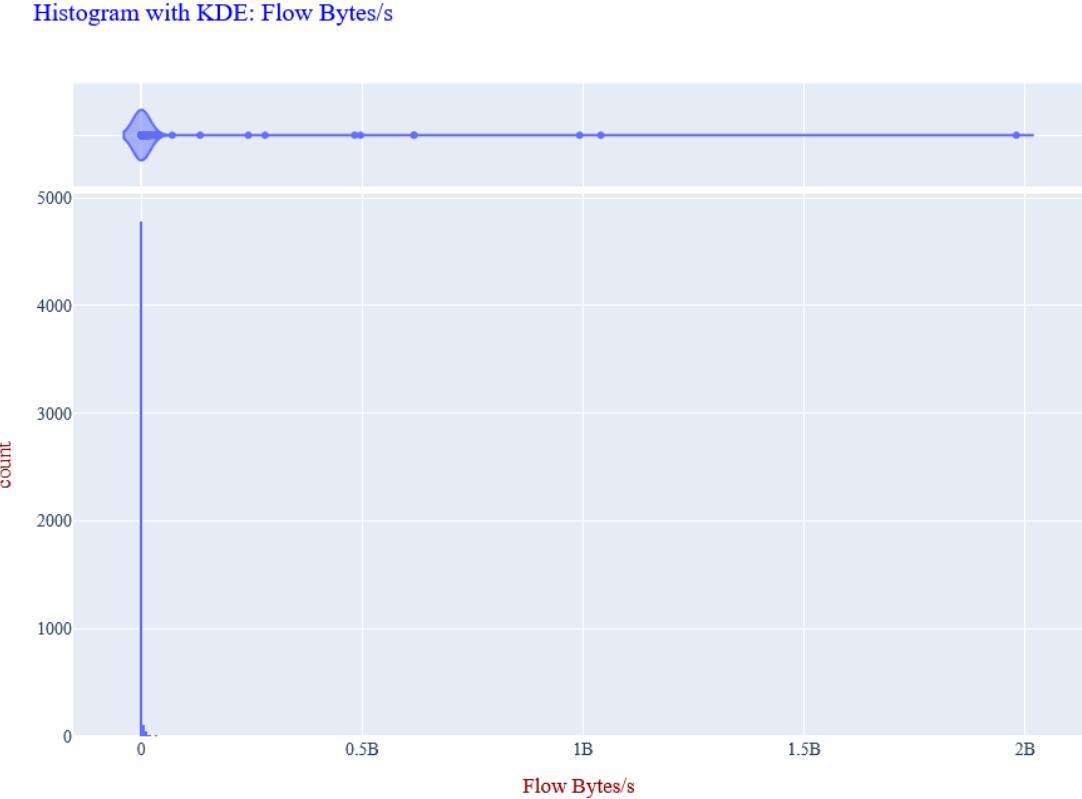


Figure 20: Histogram + KDE of raw Flow Bytes/s.

Flow Bytes/s exhibits an even more extreme concentration near zero, followed by sudden spikes into the hundreds of millions. This is typical for high-volume floods such as Hulk and GoldenEye, which produce bursts that massively outweigh benign background traffic. The KDE curve collapses toward the lower end of the domain, and the histogram bins near zero dominate the entire distribution.

Note on transformations. All histograms above intentionally use raw (untransformed) values. The dashboard includes optional `log1p`, `minmax`, and `standard` transformations, which substantially compress the heavy tails and make the structure more visible. For the purpose of documenting the true shape of the CIC-IDS2017 Wednesday file, the raw plots are more informative and help justify why transformations and outlier handling are necessary in later sections.

KDE Plot (Alpha=0.6, Width=2): Flow Duration

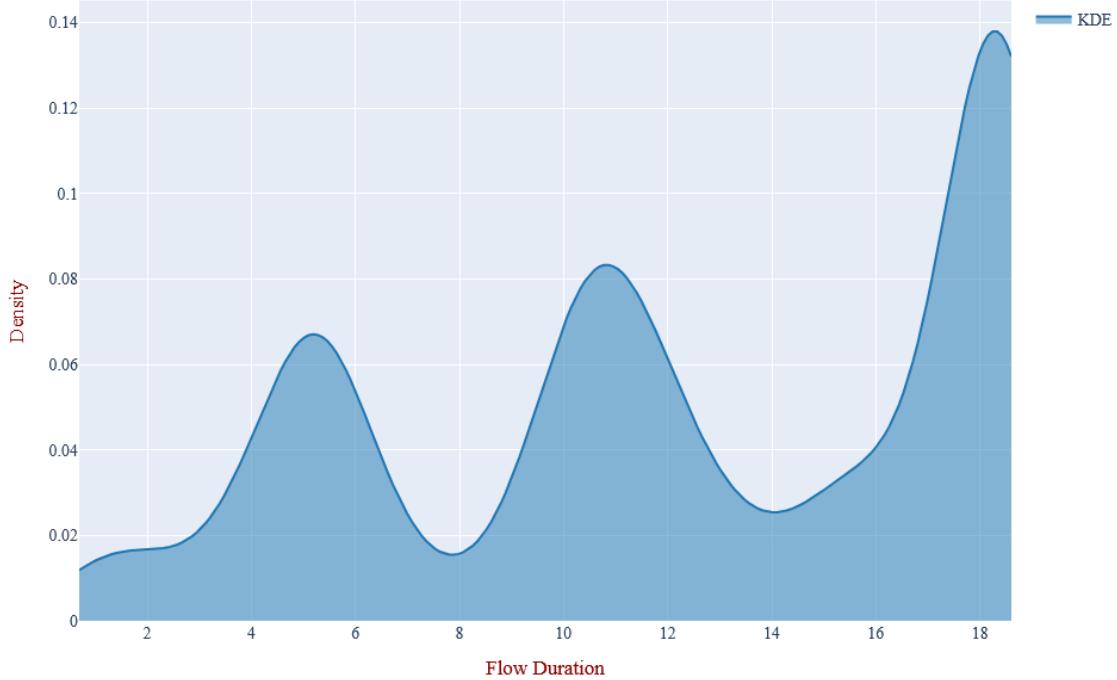


Figure 21: KDE-only visualization of Flow Duration ($\log 1p$ applied).

Observation: This KDE view removes the histogram bars and isolates the smoothed density curve. After applying $\log 1p$, the multi-modal structure becomes much easier to see. The long upper tail compresses, and the peaks line up with the attack windows observed earlier in the raw distribution. Slowloris-/SlowHTTPTest-style flows stay grouped in the long-duration low-throughput region, while Hulk-like bursts remain closer to the short-duration peak.

Rug Plot: Flow Duration

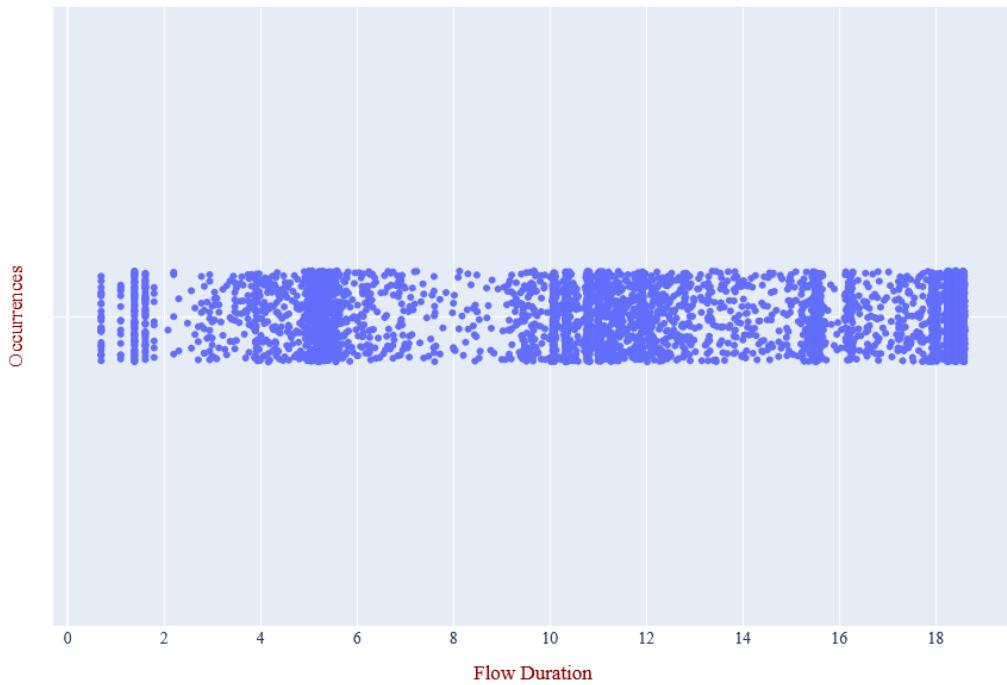


Figure 22: Rug plot for Flow Duration (log1p applied).

Observation: The rug plot shows the actual sample-level density along the axis. Even with `log1p`, flows form tight clusters that correspond to specific attack phases. This confirms that the distribution is not only heavy-tailed but also structured into distinct behavioral regimes, which is exactly what later drives the separation in PCA.

9.2 Bivariate and Multivariate Structure

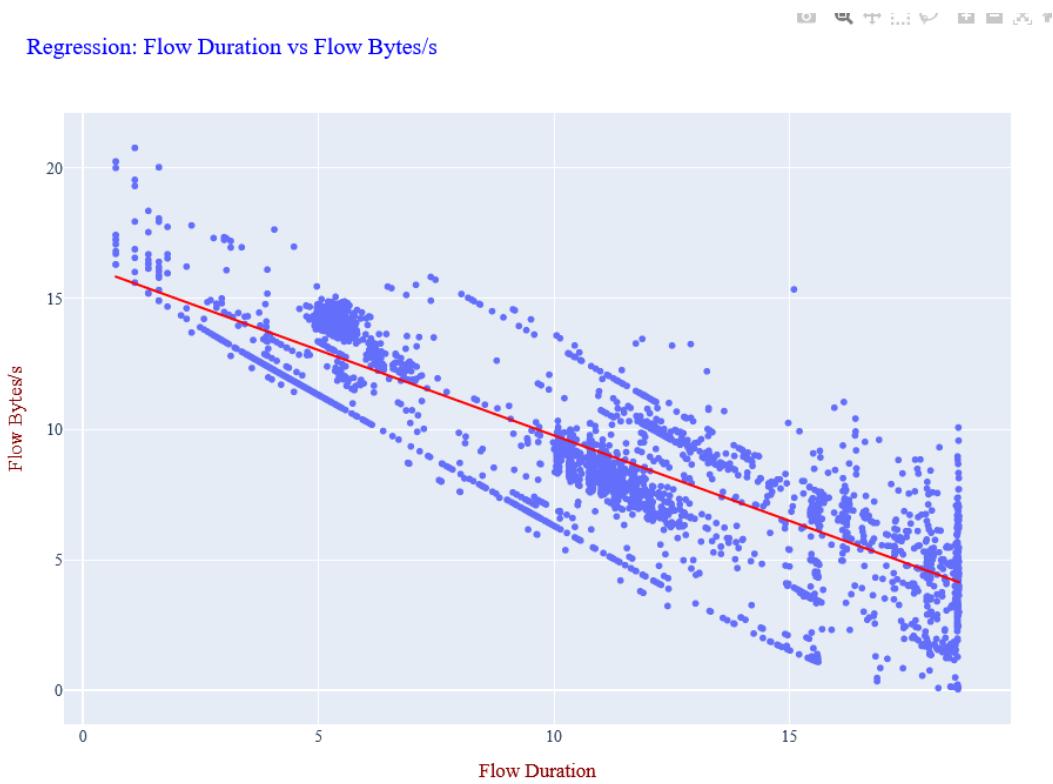


Figure 23: Regression plot of Flow Duration vs. Flow Bytes/s.

Observation: The fitted trend line slopes downward. Slowloris-style flows occupy the long-duration/low-throughput region, while Hulk floods create short, high-throughput spikes. The negative pattern is stable across subsamples and is visible even without transformations.

Contour Plot: Flow Duration vs Flow Bytes/s

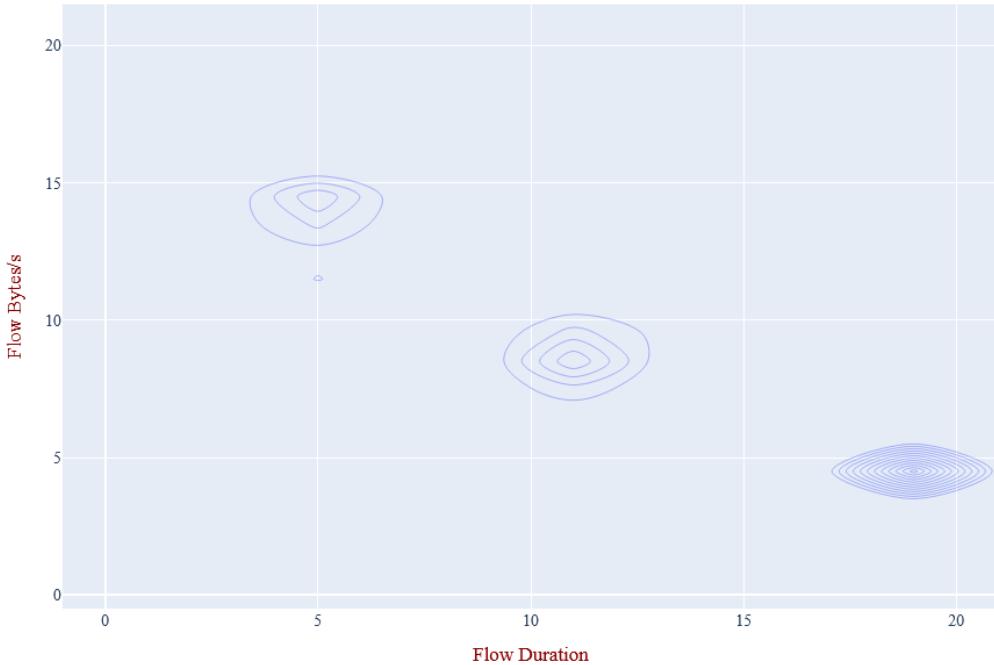


Figure 24: 2D KDE-style contour plot of Flow Duration and Flow Bytes/s.

Observation: The contour levels reveal three main density islands, corresponding to major behavior classes. Short, high-throughput flows form a compact ridge, mid-range flows cluster in a central zone, and very long, low-throughput connections remain isolated near the edge of the distribution.

Multivariate KDE (Joint Density Estimate)

The 2D KDE shown in Figure 25 is the multivariate kernel density estimate required in the statistical analysis. Instead of looking at each feature independently, this KDE models the joint density of Flow Duration and Flow Bytes/s, which are two of the strongest discriminative features in this dataset. This allows us to see how the traffic types actually cluster in the combined feature space.

Contour Plot: Flow Duration vs Flow Bytes/s

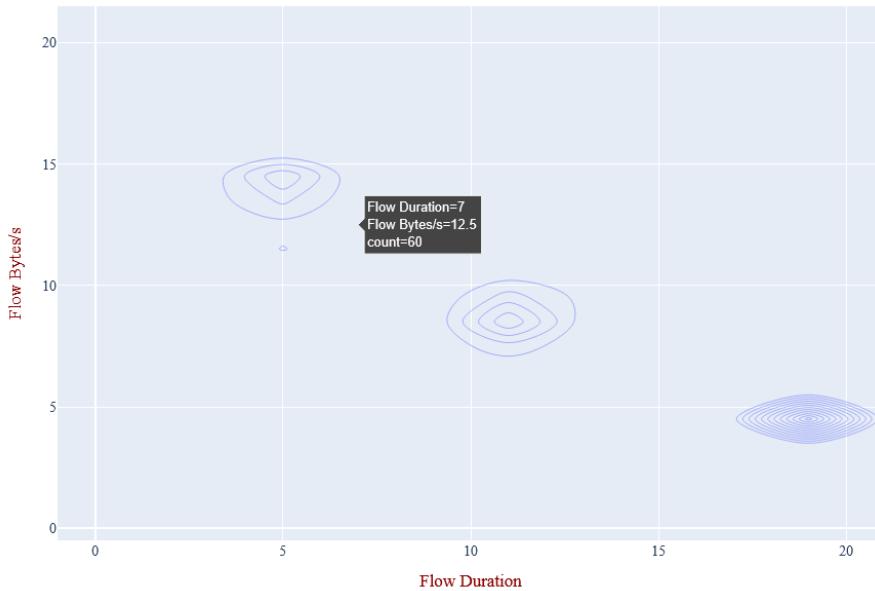


Figure 25: Multivariate (2D) KDE of Flow Duration and Flow Bytes/s.

Observation: The multivariate density surface shows that benign flows form a compact high-density region at shorter durations and moderate byte rates, while attack flows occupy entirely different areas. Hulk traffic generates a wide, elevated ridge at extremely high byte rates, whereas Slowloris and SlowHTTPTest concentrate along a low-byte, long-duration axis. These density “islands” match the separation observed in PCA and confirm that the dataset is not just non-Gaussian but also deeply multi-modal across multiple dimensions.

Hexbin Plot: Flow Duration vs Flow Bytes/s

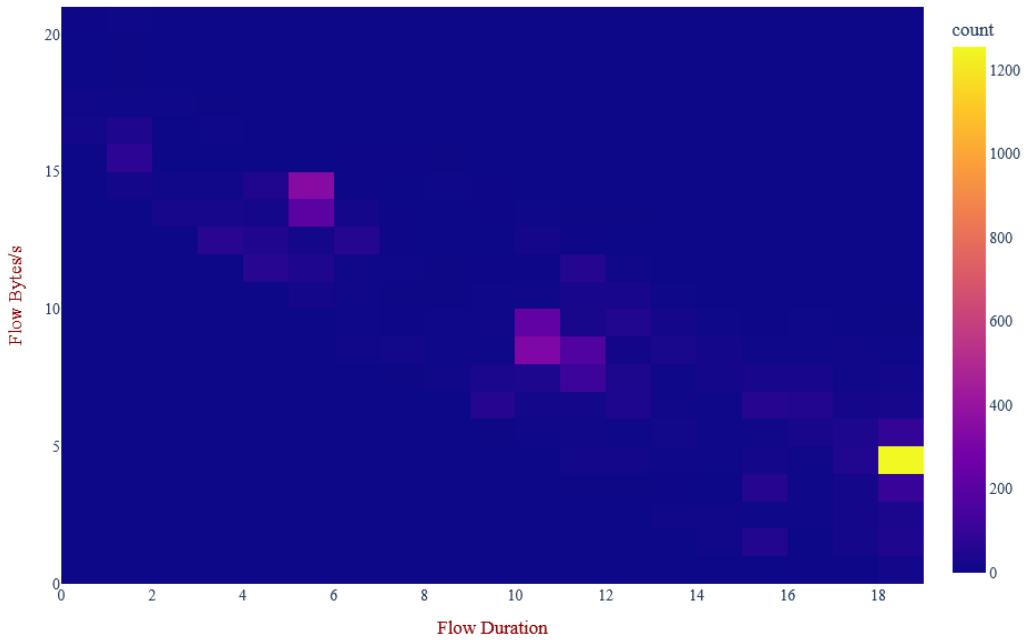


Figure 26: Hexbin density map of Flow Duration vs. Flow Bytes/s.

Observation: Hexbin counts emphasize how frequently flows accumulate in the short-duration middle-throughput band. Sparse but structured clusters appear for long-duration flows, consistent with slow-and-stuck header attacks.

Joint Plot (Scatter + KDE-style Marginals)

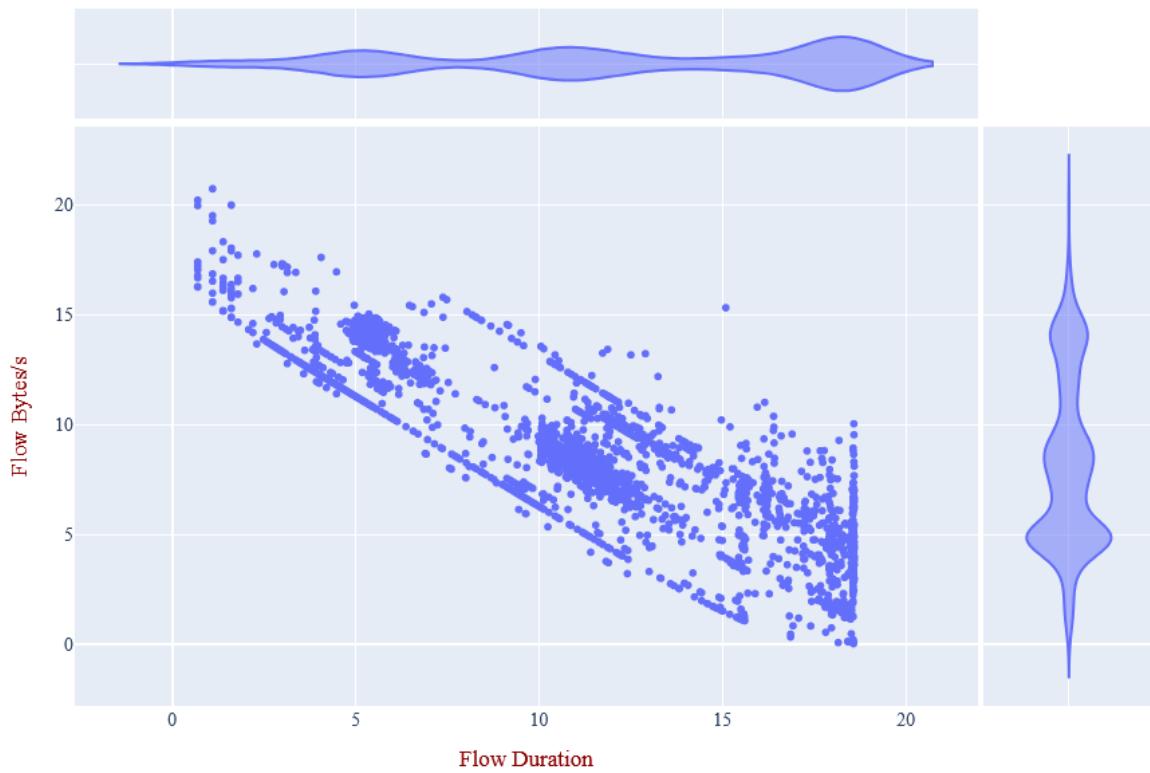


Figure 27: Joint plot with KDE marginals for Flow Duration and Flow Bytes/s.

Observation: The central scatter recreates the negative trend, while the marginal KDEs illustrate asymmetry on both axes. Duration has multiple modes, whereas byte rate is far more concentrated, especially for attack traffic.

Pair Plot

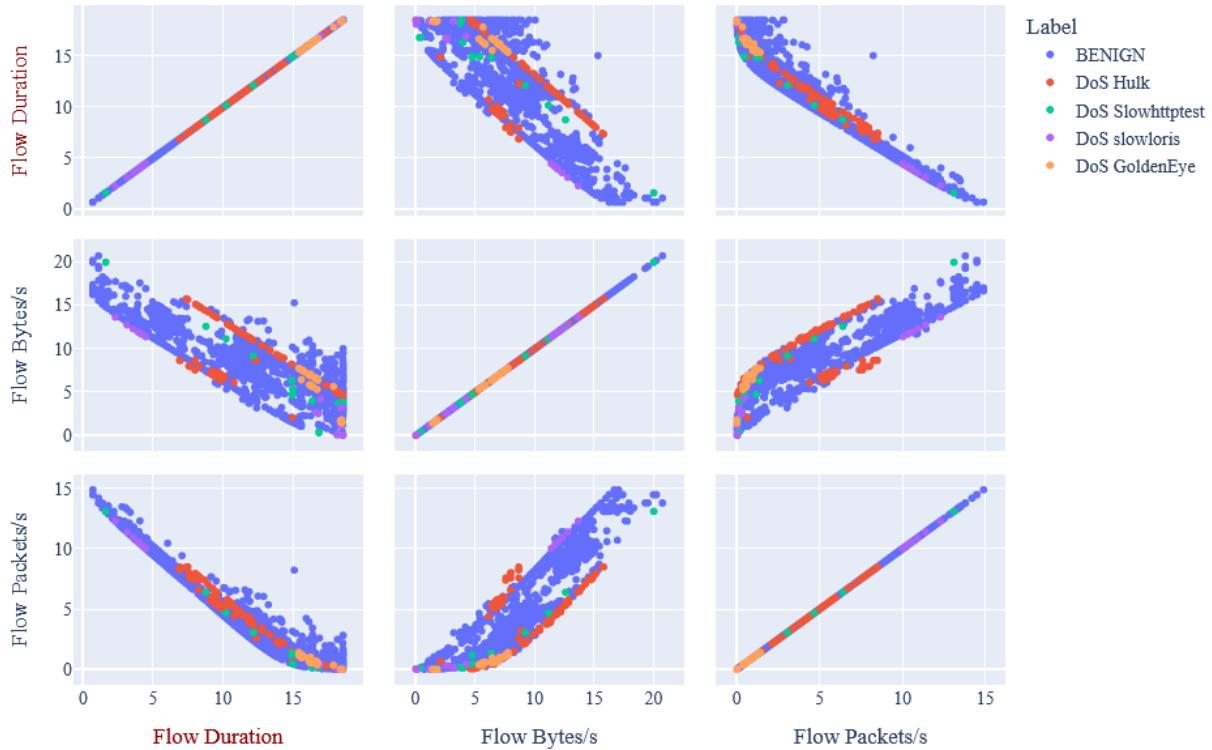


Figure 28: Pair plot for Duration, Bytes/s, Packets/s, and IAT Mean.

Observation: Attack families trace different diagonal paths in each 2D projection. Hulk shows dense, high-throughput diagonals; Slowloris and Slowhttptest anchor to long-duration, low-rate zones. Benign traffic remains dispersed, covering the central region.

3D Plot

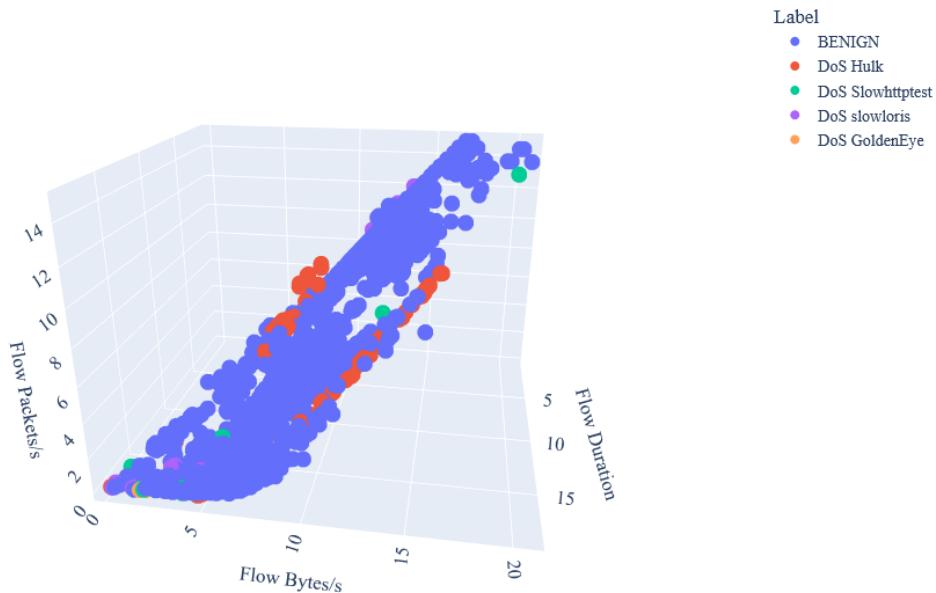


Figure 29: 3D scatter plot using Flow Duration, Flow Bytes/s, and Flow Packets/s.

Observation: In 3D space, class separation becomes even clearer. Volumetric floods (Hulk, GoldenEye) cluster in the high-packets/high-bytes corner, while header-based slow attacks occupy a stretched ridge defined by long durations but very little activity per unit time.

10 Data Visualization: Numeric and Categorical Plots

This section summarizes the static examples used to verify the dashboard outputs. The goal is to demonstrate that the dashboard not only renders interactive views but also produces consistent, interpretable plots that match offline statistical analysis.

10.1 Numeric Plots

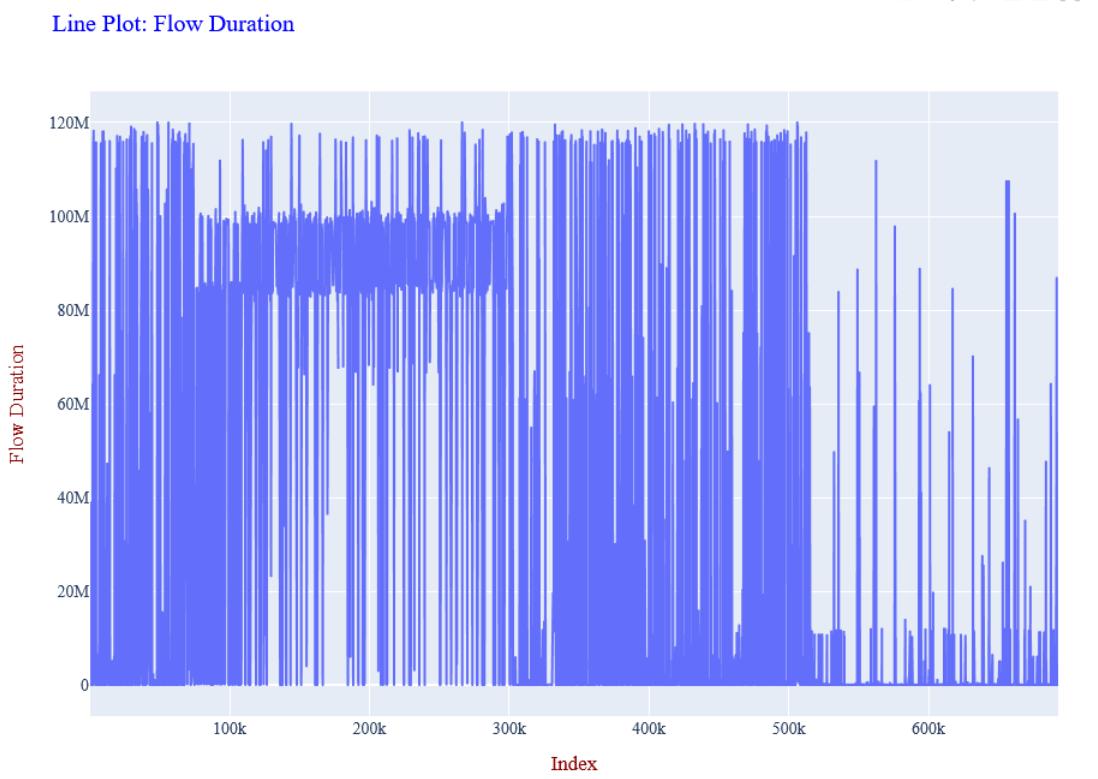


Figure 30: Line Plot of Flow Duration across row index.

Observation: The raw line plot immediately shows why this feature is so hard to interpret without preprocessing. Flow Duration jumps between extremely short and extremely long flows, and the dataset is large enough that the plot basically becomes a dense block. This reflects the mixed structure of the Wednesday file—benign flows, several DoS attacks, and large spikes during attack windows. A line plot isn't the best way to understand this feature, but it clearly shows how skewed and unstable the raw values are before applying any cleaning, outlier filtering, or transformations.

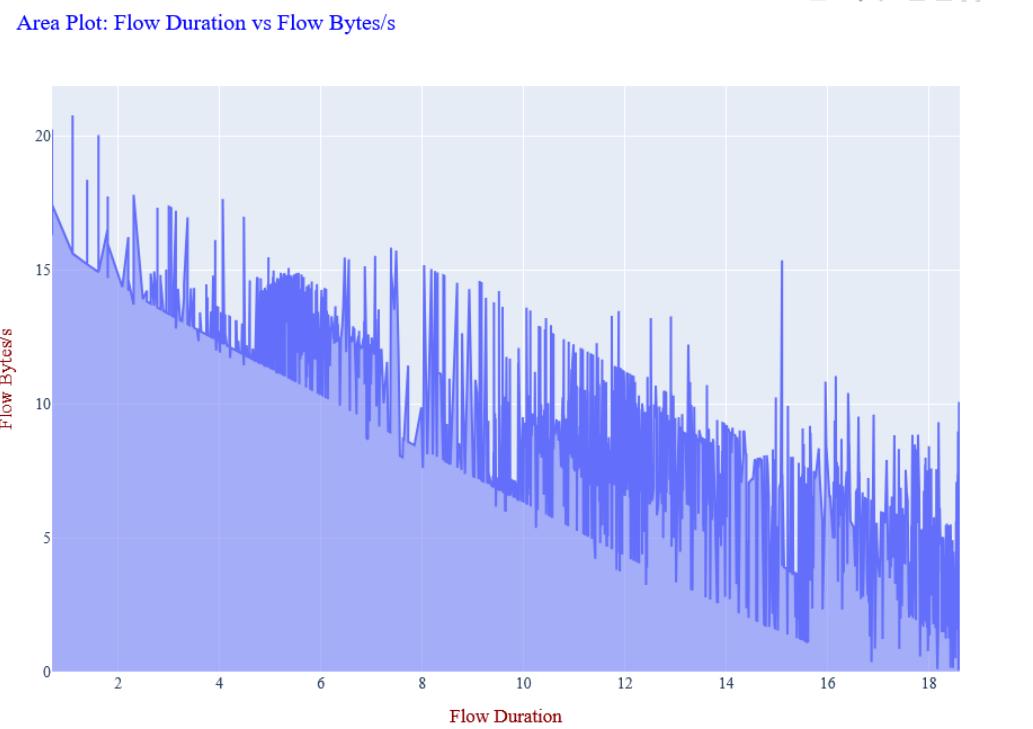


Figure 31: Area plot of attack counts across row index (time order).

Observation: The stacked area plot shows distinct blocks where DoS traffic dominates successive rows. These match the scripted attack windows inside the dataset. Outside these regions, benign traffic remains the majority.

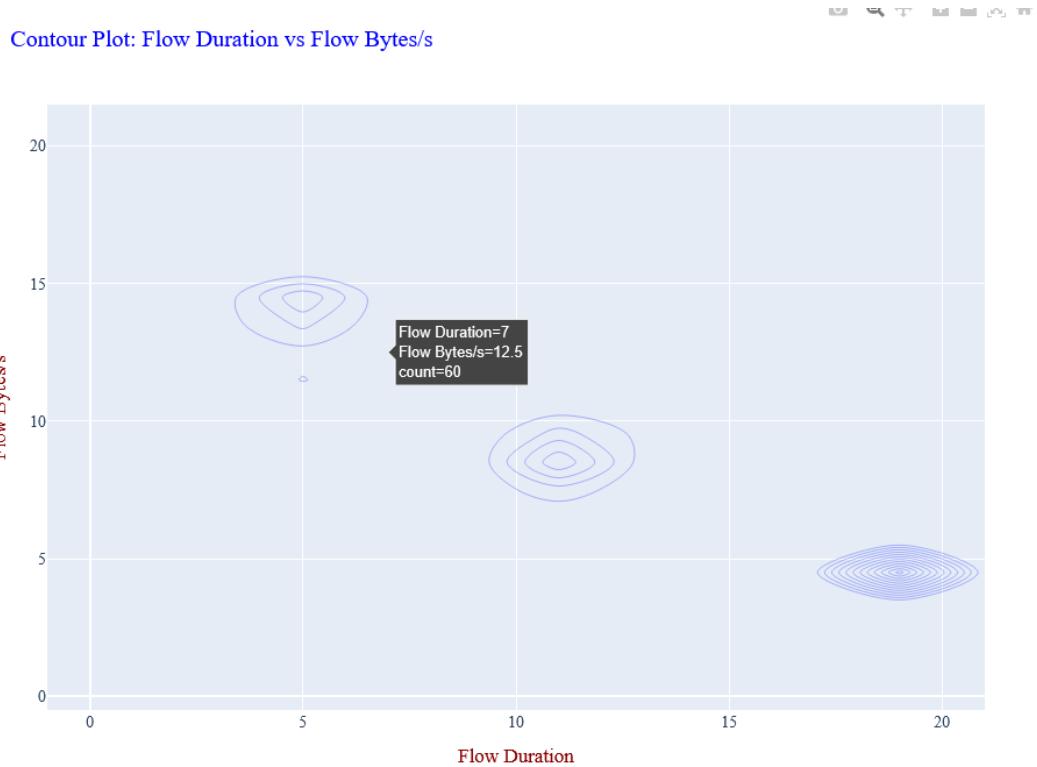


Figure 32: 2D KDE density visualization for Flow Duration and Flow Bytes/s.

Observation: The smoothed KDE surface highlights ridge structures that the hexbin plot only hints at. These ridges align with the same multi-mode behavior found in the univariate and joint plots.

Hexbin Plot: Flow Duration vs Flow Bytes/s

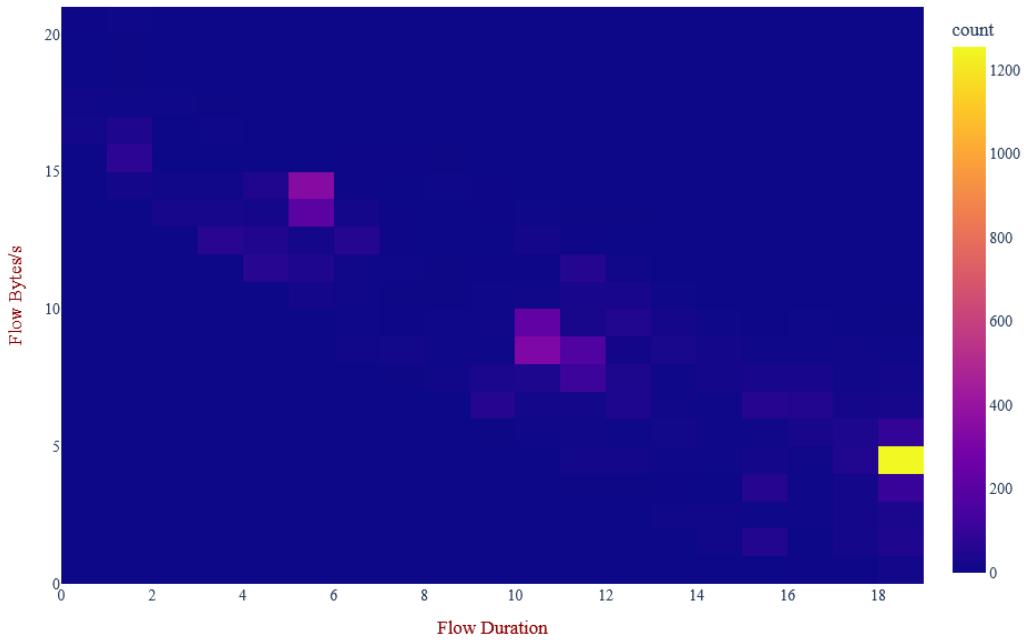


Figure 33: Hexbin visualization (numeric tab example).

Observation: The dashboard and static hexbin outputs match closely, confirming that the in-app rendering preserves the distribution geometry seen in the offline figures.

10.2 Categorical Plots

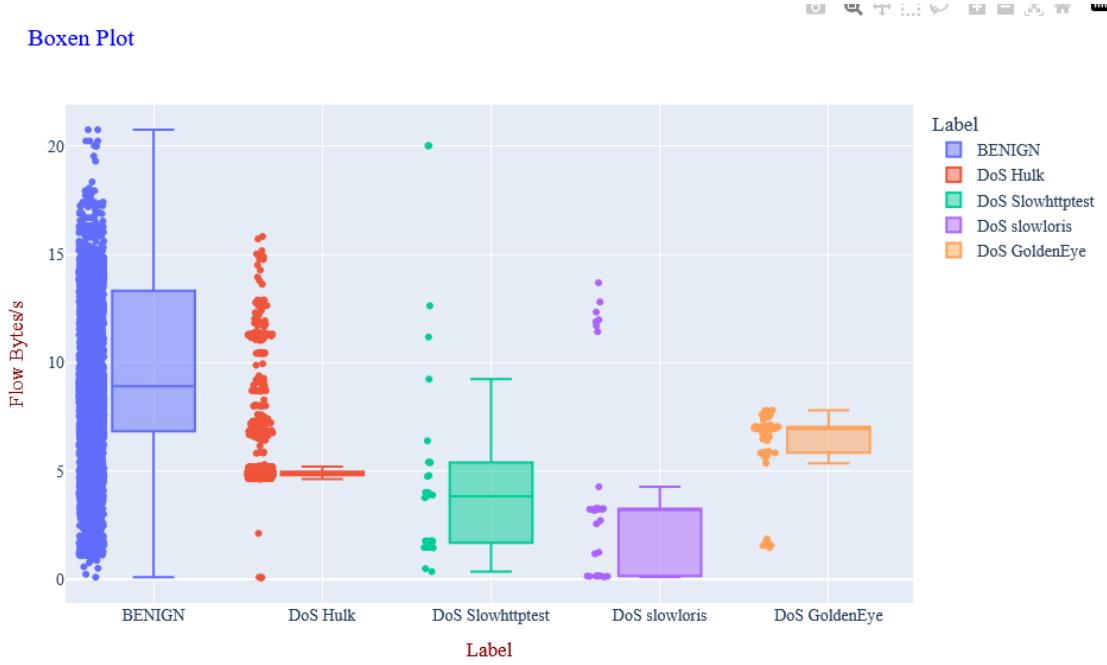


Figure 34: Boxen plot of Flow Bytes/s grouped by label.

Observation: Hulk and GoldenEye flows show compact, high-throughput bodies, while Slowloris and Slowhttptest collapse near zero. The distribution shape shifts dramatically across labels, making Flow Bytes/s one of the strongest discriminators.

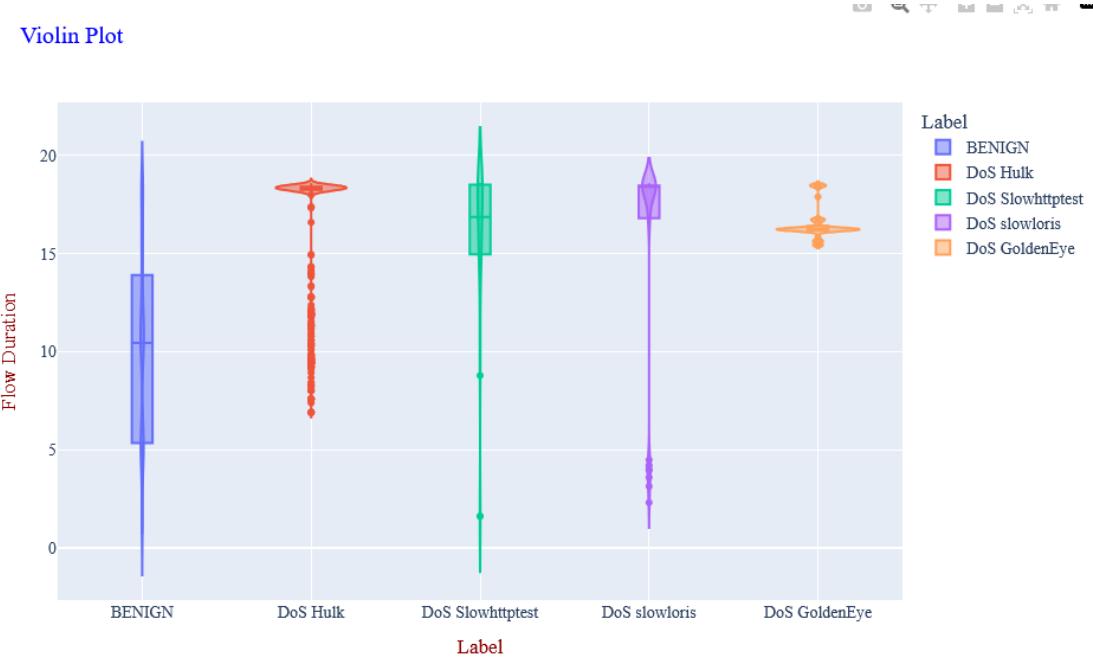


Figure 35: Violin plot of Flow Duration by label.

Observation: Slowloris and Slowhttptest exhibit narrow but very elongated duration distributions. Benign and Hulk flows appear tighter and occupy mid-range durations, mirroring the heavy-tailed nature highlighted in the statistics section.

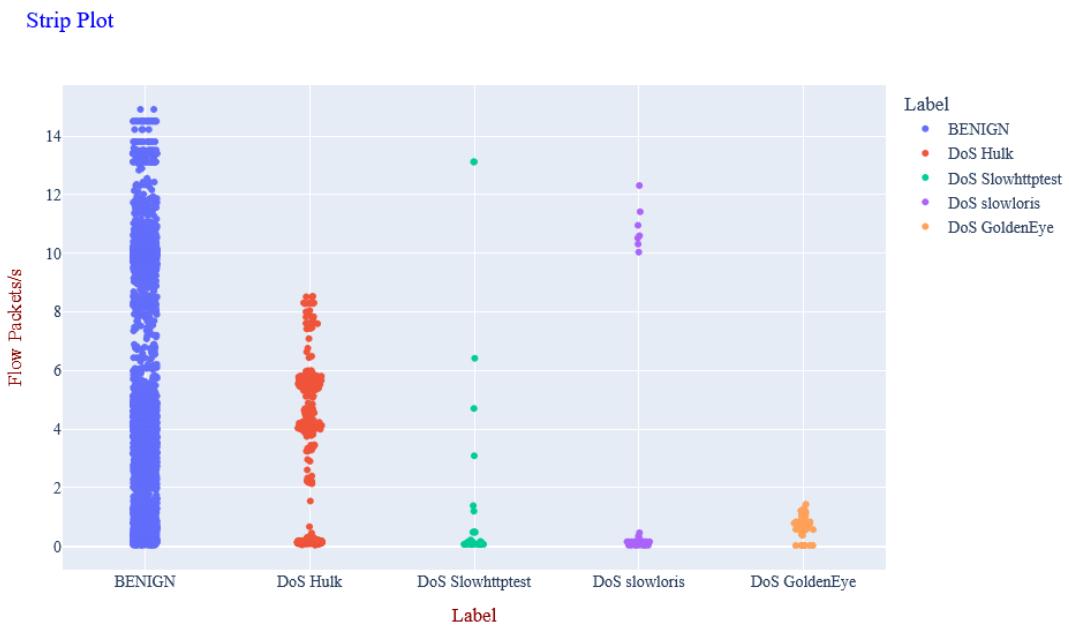


Figure 36: Strip plot of Flow Packets/s grouped by label.

Observation: The strip plot clearly separates the extremes even without bins or violins. Hulk occupies the high-packet region, Slowloris sits nearly at zero, and benign flows fill the middle band.

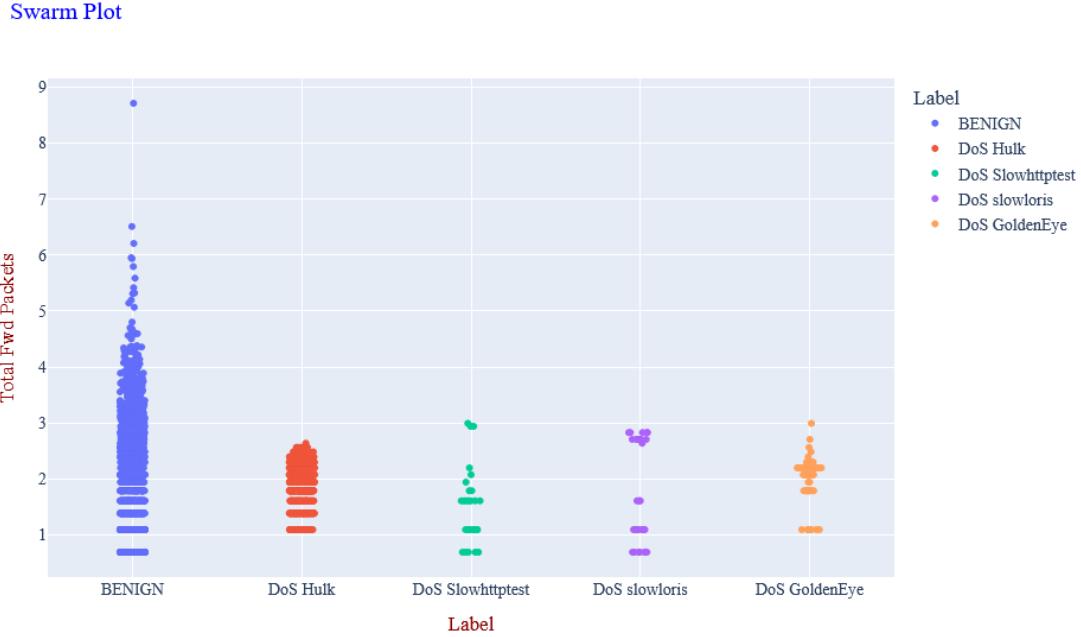


Figure 37: Swarm plot of Total Forward Packets vs. label.

Observation: The swarm layout reveals discrete packet-count bands per attack type. It avoids overplotting and makes the categorical structure more interpretable for users who prefer simple, discrete visualizations over dense histograms.

11 Subplots

The Storytelling tab uses a static four-panel subplot layout to summarize the key story of the dataset. All panels still respect the preprocessing settings selected at the top.

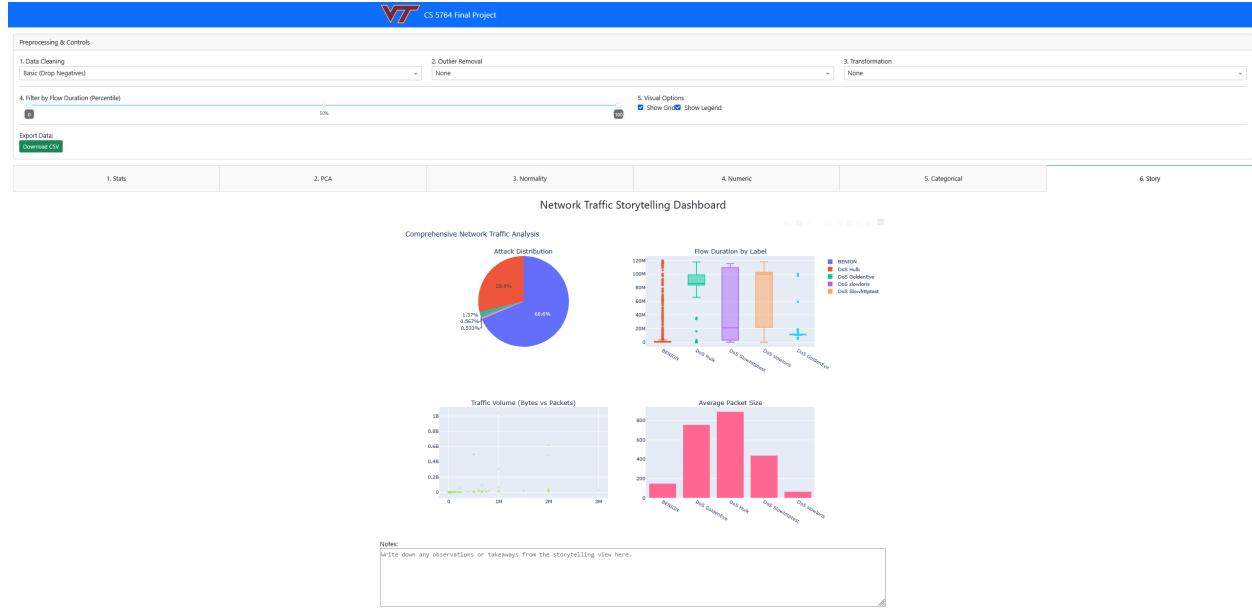


Figure 38: Storytelling dashboard tab with four linked subplots.

The layout is:

1. Top-left: Attack distribution pie chart.
2. Top-right: Box plot of Flow Duration by Label.
3. Bottom-left: Scatter of Flow Bytes vs. Flow Packets (color by label).
4. Bottom-right: Bar plot of average packet size per label.

Observation: The pie chart recaps the class imbalance. Immediately to the right, the Flow Duration box plot shows that Slowloris and Slowhttptest have far higher median durations and wider IQRs than benign traffic, which fits their behavior as connection-exhaustion attacks. The bottom-left scatter shows that Hulk and GoldenEye produce flows with huge byte and packet counts—these are volumetric floods. Finally, the average packet-size bars show that some attacks (like Hulk) favor larger packets, whereas others rely on many small packets. Put together, the four subplots give a compact visual explanation a SOC analyst could use in a meeting to explain why these attacks look different in the flow data.

12 Tables

The main tables in this project are:

- Descriptive statistics tables for numeric features (shown in the Data & Stats tab).
- The class distribution table in Section 2.
- Implicit tables within the dashboard, such as grouped means for categorical bar plots.

Observation: Compared to the plots, the tables are less visually intuitive but still necessary when I want exact values (for example, checking that the 75th percentile of Flow Duration changes after outlier removal). They also provide a more traditional view that some stakeholders prefer over graphs.

13 Dashboard

The final product of Phase III is a Plotly Dash application with the following structure:

- A global **Preprocessing & Controls** bar at the top:
 - Data Cleaning method (Keep Original / Basic / Mean Imputation / Strict).
 - Outlier Removal method (None / IQR / Z-score).
 - Transformation (None / log1p / MinMax / Standard).
 - Flow Duration percentile filter slider.
 - Visual options (grid on/off, legend on/off).
 - CSV export button to download the processed dataset.
- Six tabs:
 1. Data & Stats
 2. PCA
 3. Normality Tests
 4. Numeric Plots
 5. Categorical Plots
 6. Storytelling

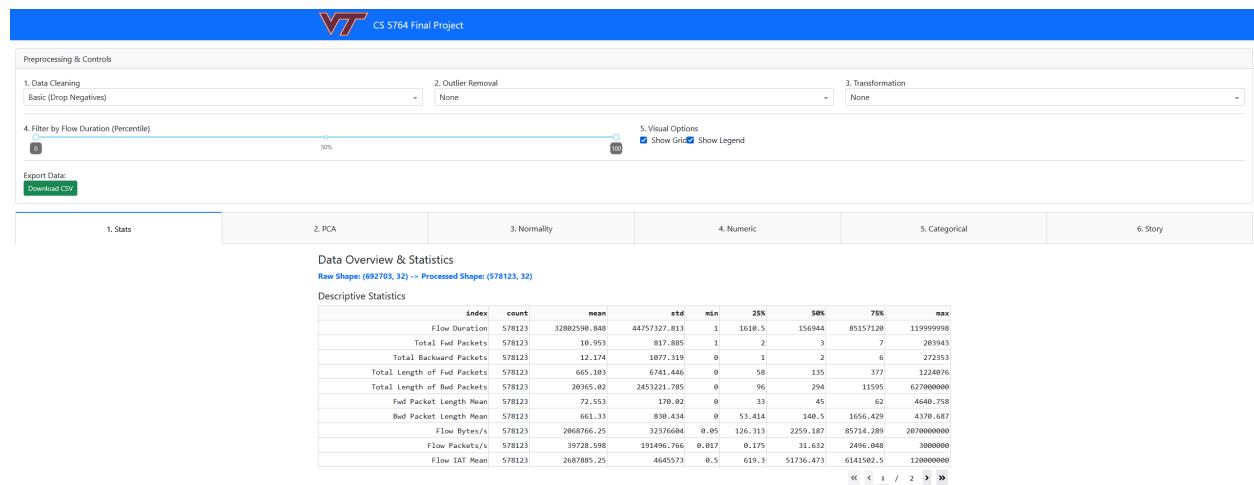


Figure 39: Dashboard main interface with preprocessing controls and Data & Stats tab.

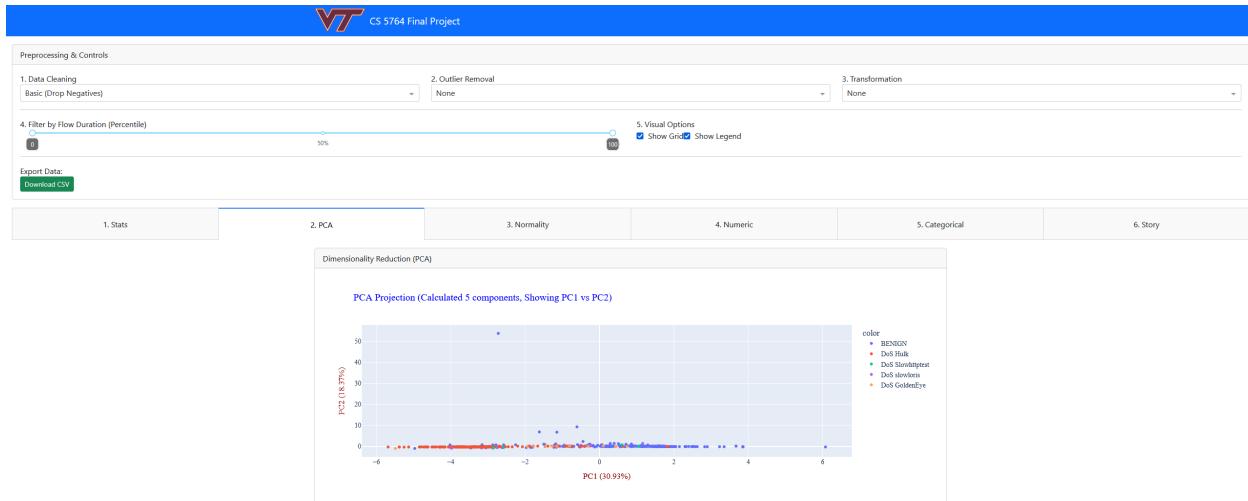


Figure 40: PCA tab showing PC1 vs. PC2 colored by label.

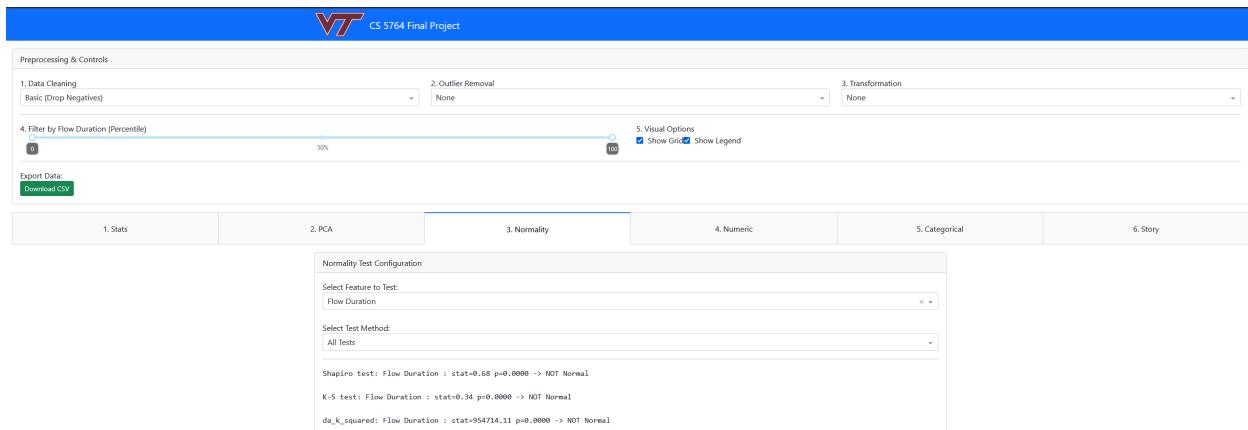


Figure 41: Normality tab: configuration and test outputs.

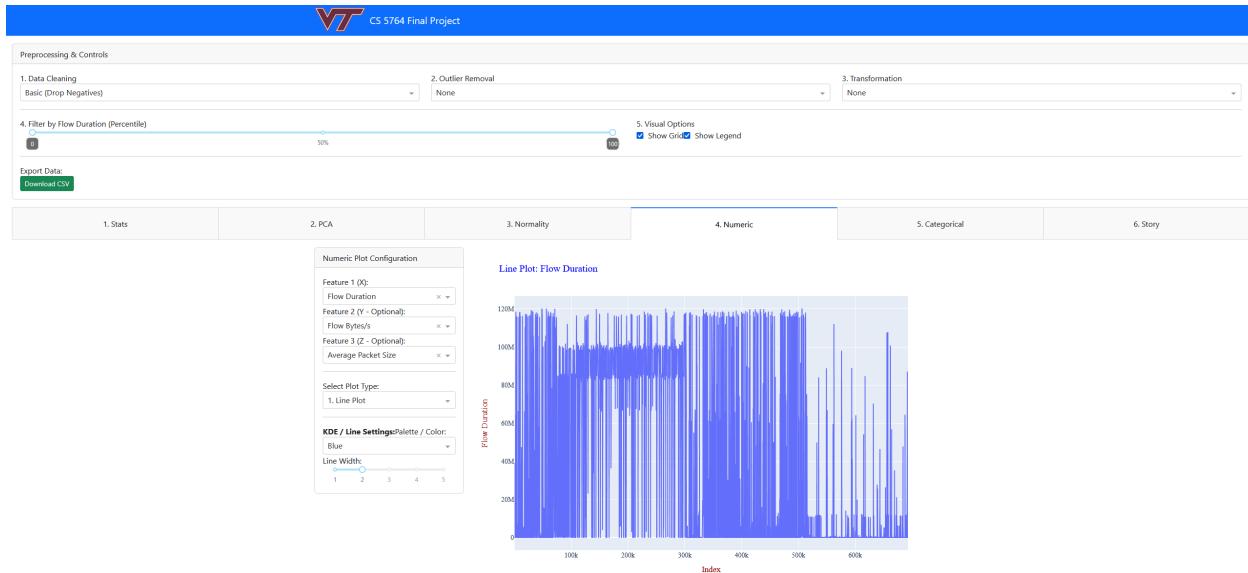


Figure 42: Numeric Plot tab.

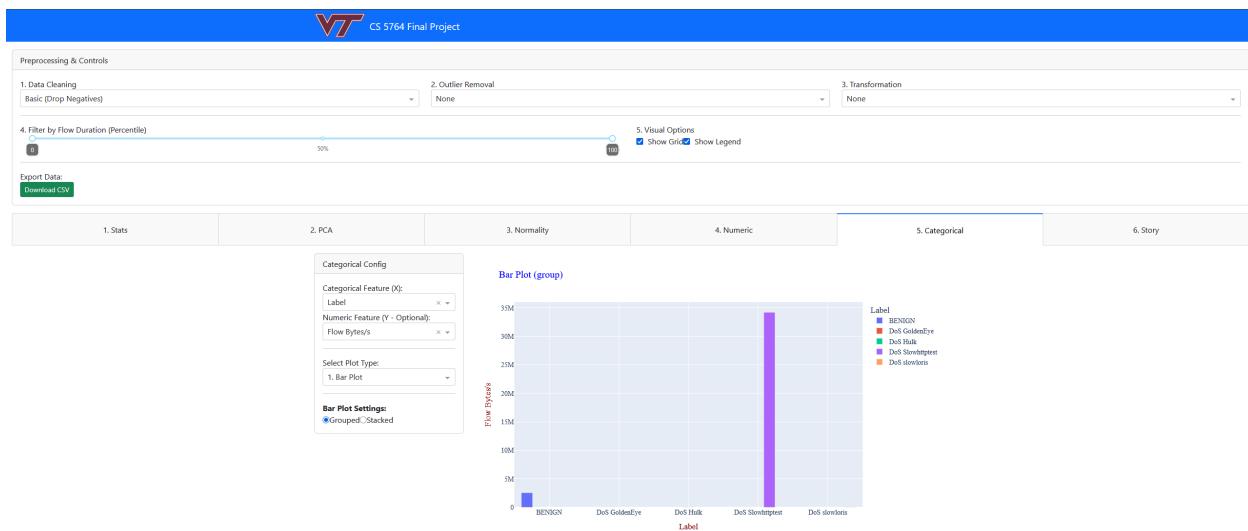


Figure 43: Categorical Plot tab.

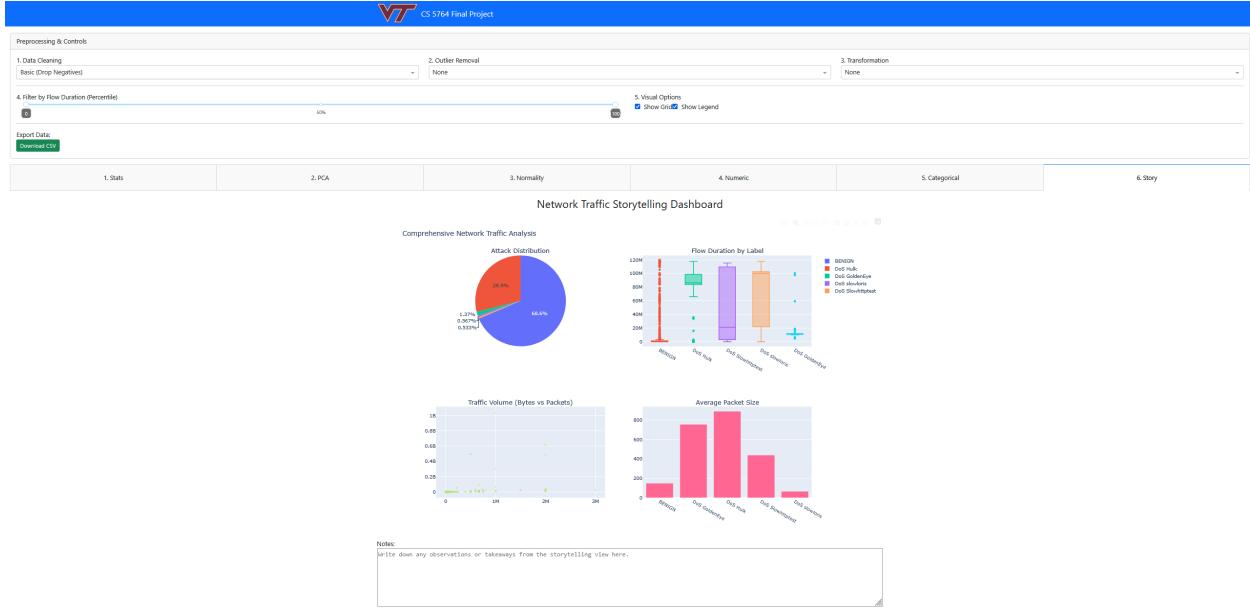


Figure 44: Storytelling tab.

Dashboard description (functionality): Every time the user touches a control—changes the cleaning method, switches from IQR to Z-score, slides the percentile filter, or picks a new feature—the corresponding callback recomputes the processed dataframethrough a single helper function and updates the visible figure. This keeps the logic consistent across tabs. The user never has to manually re-run code; the app always shows the state of the pipeline implied by the UI.

Technical Implementation (Callback Logic). The dashboard follows Dash’s callback model, meaning every control in the top panel (cleaning, outliers, transformations, percentile filters) becomes an input that forces the entire preprocessing pipeline to recompute. This is essentially a reactive setup: the dataset is always shown in the exact state implied by the UI. No cached shortcuts, no partial updates. The user interacts once, and all dependent views update immediately.

Plotly Usage. For most figures I rely on Plotly Express because it keeps the code simple and makes quick exploration easier. Wherever layout control actually matters (Storytelling tab, 3D plots, custom subplots), I switch to Graph Objects. This split keeps the code readable but still lets me build the more structured visual outputs when necessary.

14 Observations Summary

The FTP instructions emphasize that plots without observations receive zero points, so I summarize the main lessons from the entire pipeline here:

- **Class imbalance is extreme.** BENIGN and DoS Hulk dominate the dataset. Any modeling approach without rebalancing will mostly learn to distinguish “Hulk vs. everything else”.
- **Normality clearly does not hold.** Flow durations and byte/packet rates fail all three normality tests with p-values essentially zero. Heavy tails and multi-modal distributions are the norm, not the exception.
- **Preprocessing radically changes the picture.** Basic cleaning plus IQR outlier removal is enough to move the PCA from a blob dominated by outliers to a more interpretable structure where different attacks occupy different neighborhoods.
- **Correlation structure is strong.** Many features carry overlapping information. Future work could aggressively reduce dimensionality by dropping redundant columns before PCA or modeling.
- **Different DoS families have different signatures.** Slowloris/Slowhttptest focus on long-lived low-rate connections, Hulk/GoldenEye on high-volume bursts, and Heartbleed hardly appears at all. The dashboard makes these differences visible without touching any ML model.

15 Conclusion

(a) What did I learn from the graphs?

The main lesson is that network traffic is messy, skewed, and often dominated by a few behavior modes. Without cleaning, outlier control, and transformations, many standard tools (especially PCA and normality tests) produce misleading or uninterpretable results. Once I clean and transform the data, the graphs reveal clear patterns: which attacks extend flow durations, which ones saturate bandwidth, and how those behaviors interact with port usage and packet sizes.

(b) How does the dashboard help users?

Instead of hard-coding a single preprocessing pipeline, the dashboard lets the user toggle cleaning methods, outlier strategies, and transformations on the fly, then immediately see the effect on PCA, distributions, and correlation. For a SOC analyst, this is the difference between staring at static PDF plots and actually probing the dataset until it reveals something interesting. The combination of tabs covers both low-level statistics and high-level storytelling.

(c) Is the dashboard user friendly?

I shared the running dashboard with an external user (my partner) to gather informal usability feedback. The goal was to get a non-technical perspective on whether the interface, layout, and interactions communicated the intended ideas.

User Comment (verbatim):

“UI is simple, love the storytelling tab, easy to understand, but since the dataset is not familiar to me, it’s hard to understand what graph is meaning what.”

Interpretation: The feedback aligns with what I expected from a user who does not have prior exposure to intrusion–detection datasets. The storytelling tab communicates the most intuitively, which confirms its design goal. At the same time, the comment highlights that some plots require additional context when shown to non-security users. In future versions of the dashboard, adding hover–based explanations, short captions, or “why this graph matters” summaries would help bridge this gap without redesigning the plots themselves.

Overall, the feedback supports the current layout while pointing out a reasonable direction for improvement: more built-in guidance for users who are unfamiliar with CIC-IDS2017 features.

I also posted the dashboard on social media to request feedback from professionals. However, no comments were received during the period in which the post was visible. I am including this

note to document that the outreach step was completed, even though it did not result in external responses. The absence of feedback does not affect the internal evaluation of the dashboard, but it does reinforce the earlier point that users without prior exposure to intrusion-detection datasets often need additional context before engaging with the visualizations.

(d) Functionality

From a functionality standpoint, the app is not a full intrusion-detection system, but it is a solid analysis tool. It ties together every step of the FTP instructions—cleaning, outliers, PCA, normality, statistics, visualizations, and storytelling—into a single coherent pipeline. The callbacks are responsive, and the app handles a reasonably large sample of the dataset without freezing. If I wanted to evolve this into a more serious tool, the next steps would be: add time-based filtering, IP-based filtering, and integrate simple anomaly scores on top of the current visualizations.

(e) Future Work

The current dashboard stops at statistical analysis and PCA because that was the scope of the project. If I extend it later, the next step would be to integrate simple ML components directly into the same pipeline. Plotly already supports the visual side of this, so adding the following would be straightforward:

- ROC and PR curves for a basic classifier trained on the cleaned data.
- Confusion matrices to see where similar attack types get mixed.
- A lightweight anomaly detector (e.g., isolation forest) that runs on the processed dataframe and feeds into the visual tabs.

These additions would turn the dashboard from a pure analysis tool into an end-to-end exploration and modeling interface for CIC-IDS2017-style traffic.

16 Appendix: Python Code

Below is the complete Dash application source code used for this project. It includes data loading, preprocessing, outlier handling, PCA, normality tests, numeric/categorical plots, storytelling subplots, and all callbacks.

```
1 # Created by Minjin Kim on 2025.11.30
2 # CS 5764      Final Term Project
3
4 import numpy as np
5 import pandas as pd
6 import scipy.stats as stats
7 from scipy.stats import shapiro, kstest, normaltest
8 import gc
9
10 import dash
11 from dash import dcc, html, dash_table
12 from dash.dependencies import Input, Output, State
13 import dash_bootstrap_components as dbc
14
15 import plotly.express as px
16 import plotly.graph_objects as go
17 from plotly.subplots import make_subplots
18
19 from sklearn.preprocessing import StandardScaler, MinMaxScaler
20 from sklearn.decomposition import PCA
21
22 # -----
23 # 1. Global Configuration & Data Loading
24 # -----
25
26 pd.options.display.float_format = lambda x: f"{x:.2f}"
27 DATA_PATH = "Wednesday-workingHours.pcap_ISCX.csv"
28
29 NUMERIC_COLS = [
30     "Flow Duration", "Total Fwd Packets", "Total Backward Packets",
31     "Total Length of Fwd Packets", "Total Length of Bwd Packets",
32     "Fwd Packet Length Mean", "Bwd Packet Length Mean",
33     "Flow Bytes/s", "Flow Packets/s", "Flow IAT Mean",
34     "Flow IAT Std", "Fwd Packets/s", "Bwd Packets/s",
35     "Average Packet Size", "Packet Length Variance",
36     "Active Mean", "Idle Mean"
37 ]
38
```

```

39 CAT_COLS = [
40     "Destination Port", "Protocol", "Fwd PSH Flags", "Bwd PSH Flags",
41     "Fwd URG Flags", "Bwd URG Flags", "FIN Flag Count", "SYN Flag Count",
42     "RST Flag Count", "PSH Flag Count", "ACK Flag Count", "URG Flag Count",
43     "CWE Flag Count", "ECE Flag Count", "Label"
44 ]
45
46 CONST_FLAG_COLS = ["Bwd PSH Flags", "Fwd URG Flags", "Bwd URG Flags", "CWE
47     Flag Count"]
PORT_BUCKET_COL = "Destination Port Bucket"
48
49
50 def reduce_mem_usage(df):
51     start_mem = df.memory_usage().sum() / 1024 ** 2
52     print(f'Memory usage of dataframe is {start_mem:.2f} MB')
53
54     for col in df.columns:
55         col_type = df[col].dtype
56
57         if col_type != object:
58             c_min = df[col].min()
59             c_max = df[col].max()
60             if str(col_type)[:3] == 'int':
61                 if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8)
62                     .max:
63                     df[col] = df[col].astype(np.int8)
64                 elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.
65                     int16).max:
66                     df[col] = df[col].astype(np.int16)
67                 elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.
68                     int32).max:
69                     df[col] = df[col].astype(np.int32)
70                 elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.
71                     int64).max:
72                     df[col] = df[col].astype(np.int64)
73             else:
74                 df[col] = df[col].astype(np.float32)
75         else:
76             num_unique_values = len(df[col].unique())
77             num_total_values = len(df[col])
78             if num_unique_values / num_total_values < 0.5:
79                 df[col] = df[col].astype('category')
80
81

```

```

77     end_mem = df.memory_usage().sum() / 1024 ** 2
78     print(f'Memory usage after optimization is: {end_mem:.2f} MB')
79     print(f'Decreased by {100 * (start_mem - end_mem) / start_mem:.1f}%')
80
81     return df
82
83
84 def add_port_bucket_column(df: pd.DataFrame) -> pd.DataFrame:
85     if "Destination Port" not in df.columns:
86         return df
87     port = pd.to_numeric(df["Destination Port"], errors="coerce")
88     bins = [-1, 1023, 49151, 65535]
89     labels = ["Well-known (0 1023 )", "Registered (1024 49151 )", "Dynamic
90             (49152 65535 )"]
91
92     df[PORT_BUCKET_COL] = pd.cut(port, bins=bins, labels=labels)
93     df[PORT_BUCKET_COL] = df[PORT_BUCKET_COL].astype("category")
94
95     return df
96
97
98 def load_raw_dataset(path: str) -> pd.DataFrame:
99
100    df = pd.read_csv(path)
101    df.columns = df.columns.str.strip()
102    cols_present = [c for c in (NUMERIC_COLS + CAT_COLS) if c in df.columns]
103    df = df[cols_present]
104
105    df = reduce_mem_usage(df)
106
107    df = add_port_bucket_column(df)
108
109
110    print("Loading and optimizing data...")
111    df_raw = load_raw_dataset(DATA_PATH)
112    print("Data loaded successfully.")
113
114    numeric_options = [{"label": c, "value": c} for c in NUMERIC_COLS if c in
115                      df_raw.columns]
116    cat_options = [{"label": c, "value": c} for c in (CAT_COLS + [PORT_BUCKET_COL
117                ]) if c in df_raw.columns]

```

```

117         and c != "Destination Port"
118         and c not in CONST_FLAG_COLS]
119
120
121 # -----
122 # 2. Data Processing Logic
123 # -----
124
125 def clean_dataset(df: pd.DataFrame, method: str = "basic") -> pd.DataFrame:
126     df_clean = df.copy()
127
128     num_cols = df_clean.select_dtypes(include=[np.number]).columns
129     df_clean[num_cols] = df_clean[num_cols].replace([np.inf, -np.inf], np.nan)
130
131     if method in ("basic", "strict"):
132         cols_nonpositive_to_nan = ["Flow Duration", "Flow IAT Mean", "Flow
133             Bytes/s", "Flow Packets/s"]
134         for col in cols_nonpositive_to_nan:
135             if col in df_clean.columns:
136                 mask = df_clean[col] <= 0
137                 if mask.any():
138                     df_clean.loc[mask, col] = np.nan
139
140         critical = [c for c in ["Flow Duration", "Flow Bytes/s"] if c in
141             df_clean.columns]
142         if critical:
143             df_clean = df_clean.dropna(subset=critical)
144
145         if method == "strict":
146             numeric_present = [c for c in NUMERIC_COLS if c in df_clean.columns]
147             df_clean = df_clean.dropna(subset=numeric_present)
148         elif method == "fill_mean":
149             numeric_present = [c for c in NUMERIC_COLS if c in df_clean.columns]
150             df_clean[numeric_present] = df_clean[numeric_present].fillna(df_clean[
151                 numeric_present].mean())
152
153         if "Label" in df_clean.columns and df_clean["Label"].dtype == 'object':
154             df_clean["Label"] = df_clean["Label"].astype("category")
155
156         print(df_clean.head())
157         return df_clean

```

```

157 def apply_outlier_method(df: pd.DataFrame, method: str = "none") -> pd.
158     DataFrame:
159         if method == "none":
160             return df
161
162         cols = [c for c in ["Flow Duration", "Flow Bytes/s", "Average Packet Size"]
163                 if c in df.columns]
164
165         if not cols:
166             return df
167
168         # df_out = df.copy()
169         df_out = df
170
171         mask = pd.Series(True, index=df_out.index)
172
173         if method == "iqr":
174             for col in cols:
175                 q1 = df_out[col].quantile(0.25)
176                 q3 = df_out[col].quantile(0.75)
177                 iqr = q3 - q1
178                 mask &= df_out[col].between(q1 - 1.5 * iqr, q3 + 1.5 * iqr)
179
180         elif method == "zscore":
181             for col in cols:
182                 z = np.abs(stats.zscore(df_out[col].fillna(0)))
183                 mask &= (z < 3)
184
185         return df_out[mask]
186
187
188
189
190
191
192
193
194
195
196
197
def apply_transform_method(df, method="none"):
    if method == "none":
        return df

    df_tr = df.copy()
    num_cols = [c for c in NUMERIC_COLS if c in df_tr.columns]

    if method == "log1p":
        for col in num_cols:
            if df_tr[col].min() >= 0:
                df_tr[col] = np.log1p(df_tr[col])
    elif method == "minmax":
        scaler = MinMaxScaler()
        df_tr[num_cols] = scaler.fit_transform(df_tr[num_cols])

```

```

198     elif method == "standard":
199         scaler = StandardScaler()
200         df_tr[num_cols] = scaler.fit_transform(df_tr[num_cols])
201
202     return df_tr
203
204
205 def get_processed_df(clean, outlier, transform, range_val=None):
206     gc.collect()
207
208     df = clean_dataset(df_raw, method=clean)
209     df = apply_outlier_method(df, method=outlier)
210     df = apply_transform_method(df, method=transform)
211
212     if range_val and "Flow Duration" in df.columns:
213         min_p, max_p = range_val
214         if len(df) > 0:
215             low_val = np.percentile(df["Flow Duration"], min_p)
216             high_val = np.percentile(df["Flow Duration"], max_p)
217             df = df[(df["Flow Duration"] >= low_val) & (df["Flow Duration"] <=
218                     high_val)]
219
220     return df
221
222 # -----
223 # 3. Helper Functions (Including Normality Logic)
224 # -----
225
226 def make_warning_figure(msg: str) -> go.Figure:
227     fig = go.Figure()
228     fig.add_annotation(text=msg, showarrow=False, font=dict(size=16))
229     fig.update_xaxes(visible=False)
230     fig.update_yaxes(visible=False)
231     return fig
232
233
234 def perform_normality_tests(data, title, method="all"):
235     results = []
236     alpha = 0.01
237
238     sample_size = 2000
239     sample_data = data if len(data) < sample_size else np.random.choice(data,

```

```

        sample_size, replace=False)

240
241     if method in ["all", "shapiro"]:
242         stat, p = shapiro(sample_data)
243         res_str = f"Shapiro test: {title} : stat={stat:.2f} p={p:.4f}"
244         res_str += " -> Normal" if p > alpha else " -> NOT Normal"
245         results.append(res_str)

246
247     if method in ["all", "ks"]:
248         mean = np.mean(data)
249         std = np.std(data)
250         stat, p = kstest(data, 'norm', args=(mean, std))
251         res_str = f"K-S test: {title} : stat={stat:.2f} p={p:.4f}"
252         res_str += " -> Normal" if p > alpha else " -> NOT Normal"
253         results.append(res_str)

254
255     if method in ["all", "k2"]:
256         try:
257             stat, p = normaltest(data)
258             res_str = f"da_k_squared: {title} : stat={stat:.2f} p={p:.4f}"
259             res_str += " -> Normal" if p > alpha else " -> NOT Normal"
260             results.append(res_str)
261         except Exception as e:
262             results.append(f"da_k_squared failed: {str(e)}")

263
264     return "\n\n".join(results)

265
266
267 # -----
268 # 4. Layout & App
269 # -----
270 external_stylesheets = [dbc.themes.BOOTSTRAP]
271 app = dash.Dash(__name__, external_stylesheets=external_stylesheets)
272 server = app.server

273
274 navbar = dbc.Navbar(
275     dbc.Container([
276         html.A(
277             dbc.Row([
278                 dbc.Col(
279                     html.Img(src="https://upload.wikimedia.org/wikipedia/
280                         commons/6/60/Virginia_Tech_Hokies_logo.svg",
281                         height="50px"))),

```

```

281         dbc.Col(dbc.NavbarBrand("CS 5764 Final Project", className="ms-2"),
282                     align="center", className="g-0"),
283                     href="#", style={"textDecoration": "none"}, ,
284                     ),
285     ],
286     color="primary", dark=True,
287 )
288
289 # --- Global Controls ---
290 controls = dbc.Card([
291     dbc.CardHeader("Preprocessing & Controls"),
292     dbcCardBody([
293         dbc.Row([
294             dbc.Col([
295                 html.Label("1. Data Cleaning"),
296                 dcc.Dropdown(
297                     id="clean-method",
298                     options=[
299                         {"label": "Keep Original", "value": "none"},
300                         {"label": "Basic (Drop Negatives)", "value": "basic"},
301                         {"label": "Strict (Drop All NaNs)", "value": "strict"}
302                         },
303                         {"label": "Fill Mean", "value": "fill_mean"}
304                     ],
305                     value="basic", clearable=False
306                 )
307             ],
308             width=4),
309             dbc.Col([
310                 html.Label("2. Outlier Removal"),
311                 dcc.Dropdown(
312                     id="outlier-method",
313                     options=[
314                         {"label": "None", "value": "none"},
315                         {"label": "IQR Method", "value": "iqr"},
316                         {"label": "Z-Score Method", "value": "zscore"}
317                     ],
318                     value="none", clearable=False
319                 )
320             ],
321             width=4),
322             dbc.Col([
323                 html.Label("3. Transformation"),
324                 dcc.Dropdown(

```

```

322         id="transform-method",
323         options=[
324             {"label": "None", "value": "none"},
325             {"label": "Log1p", "value": "log1p"},
326             {"label": "MinMax Scaling", "value": "minmax"},
327             {"label": "Standard Scaling", "value": "standard"}
328         ],
329         value="none", clearable=False
330     )
331     ], width=4),
332 ],
333 html.Hr(),
334 dbc.Row([
335     dbc.Col([
336         html.Label("4. Filter by Flow Duration (Percentile)"),
337         dcc.RangeSlider(
338             id='range-slider',
339             min=0, max=100, step=5,
340             value=[0, 100],
341             marks={0: '0%', 50: '50%', 100: '100%'},
342             tooltip={"placement": "bottom", "always_visible": True}
343         )
344     ], width=6),
345     dbc.Col([
346         html.Label("5. Visual Options"),
347         dcc.Checklist(
348             id='graph-options',
349             options=[
350                 {'label': ' Show Grid', 'value': 'grid'},
351                 {'label': ' Show Legend', 'value': 'legend'}
352             ],
353             value=['grid', 'legend'],
354             inline=True,
355             inputStyle={"margin-right": "5px"}
356         )
357     ], width=6)
358 ],
359 html.Hr(),
360 dbc.Row([
361     dbc.Col([
362         html.Label("Export Data:"),
363         html.Br(),
364         dbc.Button("Download CSV", id="btn-download", color="success",

```

```

            size="sm"),
365        dcc.Download(id="download-dataframe-csv"),
366        dbc.Tooltip(
367            "Download the currently processed dataset as CSV",
368            target="btn-download",
369        )
370    ], width=12)
371 )
372 ]
373 ], className="mb-3")

374
375 # --- Tab 1: Data & Stats ---
376 tab1 = dbc.Container([
377     html.H4("Data Overview & Statistics"),
378     html.Div(id="data-shape-info", className="mb-3 text-primary fw-bold"),
379     html.H5("Descriptive Statistics"),
380     html.Div(id="stats-table-container")
381 ], className="mt-3")

382
383 # --- Tab 2: PCA ---
384 tab2 = dbc.Container([
385     dbc.Row([
386         dbc.Col([
387             dbc.Card([
388                 dbc.CardHeader("Dimensionality Reduction (PCA)"),
389                 dbcCardBody(
390                     dcc.Loading(
391                         id="loading-pca",
392                         type="default",
393                         children=dcc.Graph(id="pca-graph")
394                     )
395                 )
396             ],
397             width=12),
398         ])
399     ], className="mt-3"))

400
401 # --- Tab 3: Normality Tests ---
402 tab3 = dbc.Container([
403     dbc.Row([
404         dbc.Col([
405             dbc.Card([
406                 dbc.CardHeader("Normality Test Configuration"),

```

```

407     dbc.CardBody([
408         html.Label("Select Feature to Test:"),
409         dcc.Dropdown(id="norm-feature", options=numeric_options,
410                     value="Flow Duration"),
411         html.Br(),
412         html.Label("Select Test Method:"),
413         dcc.Dropdown(
414             id="norm-method",
415             options=[
416                 {'label': 'All Tests', 'value': 'all'},
417                 {'label': 'Shapiro-Wilk', 'value': 'shapiro'},
418                 {'label': 'Kolmogorov-Smirnov (K-S)', 'value': 'ks',
419                  },
420                 {'label': "D'Agostino's K-squared", 'value': 'k2'}
421             ],
422             value='all', clearable=False
423         ),
424         html.Hr(),
425         html.Div(id="norm-result-text", style={"whiteSpace": "pre-
426             wrap", "fontFamily": "monospace"})
427     ])
428 ],
429 className="mt-3")
430
431 # --- Tab 4: Numeric Plots ---
432 tab4 = dbc.Container([
433     dbc.Row([
434         dbc.Col([
435             dbc.Card([
436                 dbc.CardHeader("Numeric Plot Configuration"),
437                 dbc.CardBody([
438                     html.Label("Feature 1 (X):"),
439                     dcc.Dropdown(id="num-f1", options=numeric_options, value="Flow Duration"),
440                     html.Label("Feature 2 (Y - Optional):"),
441                     dcc.Dropdown(id="num-f2", options=numeric_options, value="Flow Bytes/s"),
442                     html.Label("Feature 3 (Z - Optional):"),
443                     dcc.Dropdown(id="num-f3", options=numeric_options, value="Average Packet Size"),
444                     html.Hr(),
445                 ])
446             ])
447         ])
448     ])
449 ])

```

```

444     html.Label("Select Plot Type:"),  

445     dcc.Dropdown(  

446         id="num-type",  

447         options=[  

448             {"label": "1. Line Plot", "value": "line"},  

449             {"label": "2. Dist Plot (Histogram)", "value": "dist"},  

450             {"label": "3. Pair Plot", "value": "pair"},  

451             {"label": "4. Heatmap (Correlation)", "value": "heatmap"},  

452             {"label": "5. Histogram + KDE", "value": "hist_kde"},  

453             {"label": "6. QQ Plot", "value": "qq"},  

454             {"label": "7. KDE Plot (Custom)", "value": "kde_custom"},  

455             {"label": "8. Lm/Reg Plot", "value": "reg"},  

456             {"label": "9. Area Plot", "value": "area"},  

457             {"label": "10. Joint Plot", "value": "joint"},  

458             {"label": "11. Rug Plot", "value": "rug"},  

459             {"label": "12. 3D Plot", "value": "3d"},  

460             {"label": "13. Contour Plot", "value": "contour"},  

461             {"label": "14. Cluster Map", "value": "cluster"},  

462             {"label": "15. Hexbin Plot", "value": "hexbin"}  

463         ],  

464         value="line", clearable=False  

465     ),  

466     html.Hr(),  

467     html.Label("KDE / Line Settings:", className="fw-bold"),  

468     html.Label("Palette / Color:"),  

469     dcc.Dropdown(  

470         id="kde-palette",  

471         options=[  

472             {"label": "Blue", "value": "blue"},  

473             {"label": "Red", "value": "red"},  

474             {"label": "Green", "value": "green"},  

475             {"label": "Purple", "value": "purple"}  

476         ],  

477         value="blue", clearable=False  

478     ),  

479     html.Label("Line Width:"),  

480     dcc.Slider(id="kde-width", min=1, max=5, step=0.5, value  

481         =2,  

482         marks={1: '1', 2: '2', 3: '3', 4: '4', 5: '5'})  


```

```

482         ] )
483     ])
484   ], width=3),
485   dbc.Col([
486     dcc.Loading(
487       id="loading-num",
488       type="default",
489       children=dcc.Graph(id="numeric-graph", style={"height": "700px"
490         })
491     ], width=9)
492   ])
493 ], className="mt-3")
494
495 # --- Tab 5: Categorical Plots ---
496 tab5 = dbc.Container([
497   dbc.Row([
498     dbc.Col([
499       dbc.Card([
500         dbc.CardHeader("Categorical Config"),
501         dbc.CardBody([
502           html.Label("Categorical Feature (X):"),
503           dcc.Dropdown(id="cat-col", options=cat_options, value="Label"),
504           html.Label("Numeric Feature (Y - Optional):"),
505           dcc.Dropdown(id="cat-num", options=numeric_options, value="Flow Bytes/s"),
506           html.Hr(),
507           html.Label("Select Plot Type:"),  

508           dcc.Dropdown(
509             id="cat-type",
510             options=[
511               {"label": "1. Bar Plot", "value": "bar"},  

512               {"label": "2. Count Plot", "value": "count"},  

513               {"label": "3. Pie Chart", "value": "pie"},  

514               {"label": "4. Multivariate Box", "value": "box_multi"},  

515               {"label": "5. Multivariate Boxen", "value": "boxen_multi"},  

516               {"label": "6. Violin Plot", "value": "violin"},  

517               {"label": "7. Strip Plot", "value": "strip"},  

518               {"label": "8. Swarm Plot", "value": "swarm"}  

519             ],
520           ]
521         )
522       ])
523     ])
524   ])
525 ]
526 
```

```

520             value="bar", clearable=False
521         ),
522         html.Hr(),
523         html.Label("Bar Plot Settings:", className="fw-bold"),
524         dcc.RadioItems(
525             id="bar-mode",
526             options=[
527                 {"label": "Grouped", "value": "group"},
528                 {"label": "Stacked", "value": "stack"}
529             ],
530             value="group",
531             inline=True
532         )
533     ]
534 )
535 ], width=3),
536 dbc.Col([
537     dcc.Loading(
538         id="loading-cat",
539         type="default",
540         children=dcc.Graph(id="cat-graph", style={"height": "600px"})
541     )
542 ], width=9)
543 ])
544 ], className="mt-3")
545
546 # --- Tab 6: Storytelling ---
547 tab6 = dbc.Container([
548     html.H3("Network Traffic Storytelling Dashboard", className="text-center
549             my-3"),
550     dcc.Loading(
551         id="loading-story",
552         type="default",
553         children=dcc.Graph(id="story-graph", style={"height": "800px"})
554     ),
555
556     html.Br(),
557     html.Label("Notes:"),
558     dcc.Textarea(
559         id="story-notes",
560         placeholder="Write down any observations or takeaways from the
storytelling view here."

```

```

561         style={
562             "width": "100%",
563             "height": "150px",
564             "fontFamily": "monospace"
565         }
566     )
567 ], className="mt-3")
568
569 # --- Main Layout ---
570 app.layout = html.Div([
571     navbar,
572     dbc.Container([
573         controls,
574         dcc.Tabs(
575             id="tabs",
576             value="tab-1",
577             children=[
578                 dcc.Tab(label="1. Stats", value="tab-1", children=tab1),
579                 dcc.Tab(label="2. PCA", value="tab-2", children=tab2),
580                 dcc.Tab(label="3. Normality", value="tab-3", children=tab3),
581                 dcc.Tab(label="4. Numeric", value="tab-4", children=tab4),
582                 dcc.Tab(label="5. Categorical", value="tab-5", children=tab5),
583                 dcc.Tab(label="6. Story", value="tab-6", children=tab6),
584             ],
585         ),
586     ], fluid=True, className="mt-3"),
587 ])
588
589
590 # -----
591 # 5. Callbacks
592 # -----
593
594 @app.callback(
595     Output("download-dataframe-csv", "data"),
596     Input("btn-download", "n_clicks"),
597     State("clean-method", "value"),
598     State("outlier-method", "value"),
599     State("transform-method", "value"),
600     State("range-slider", "value"),
601     prevent_initial_call=True
602 )
603 def download_processed_data(n_clicks, clean, outlier, transform, range_val):

```

```

604     if n_clicks is None:
605         return dash.no_update
606     df = get_processed_df(clean, outlier, transform, range_val)
607     return dcc.send_data_frame(df.to_csv, "processed_data.csv")
608
609
610 @app.callback(
611     Output("data-shape-info", "children"),
612     Output("stats-table-container", "children"),
613     Input("clean-method", "value"),
614     Input("outlier-method", "value"),
615     Input("transform-method", "value"),
616     Input("range-slider", "value")
617 )
618 def update_stats(clean, outlier, transform, range_val):
619     df = get_processed_df(clean, outlier, transform, range_val)
620     msg = f"Raw Shape: {df_raw.shape} -> Processed Shape: {df.shape}"
621
622     desc = df[NUMERIC_COLS].describe().T.reset_index()
623     desc = desc.round(3)
624
625     table = dash_table.DataTable(
626         data=desc.to_dict('records'),
627         columns=[{"name": i, "id": i} for i in desc.columns],
628         style_table={'overflowX': 'auto'},
629         page_size=10,
630         style_header={'fontWeight': 'bold'}
631     )
632     return msg, table
633
634
635 @app.callback(
636     Output("pca-graph", "figure"),
637     Input("clean-method", "value"),
638     Input("outlier-method", "value"),
639     Input("transform-method", "value"),
640     Input("range-slider", "value")
641 )
642 def update_pca(clean, outlier, transform, range_val):
643     df = get_processed_df(clean, outlier, transform, range_val)
644     if df.empty:
645         return make_warning_figure("No Data after filtering")
646

```

```

647 df_sub = df.sample(min(len(df), 1000), random_state=42)
648
649 X = df_sub[NUMERIC_COLS].select_dtypes(include=[np.number]).fillna(0)
650
651 n_samples, n_features = X.shape
652
653 if n_samples < 2 or n_features < 2:
654     return make_warning_figure("Not enough data for PCA")
655
656 target_n = 5
657 n_components = min(target_n, n_samples, n_features)
658
659 X_std = StandardScaler().fit_transform(X)
660
661 pca = PCA(n_components=n_components)
662 components = pca.fit_transform(X_std)
663
664 # print(f"PCA performed with n_components={n_components}")
665
666 color_col = df_sub["Label"] if "Label" in df_sub.columns else None
667 if color_col is not None and hasattr(color_col, 'cat'):
668     color_col = color_col.astype(str)
669
670 pca_fig = px.scatter(
671     x=components[:, 0],
672     y=components[:, 1],
673     color=color_col,
674     title=f"PCA Projection (Calculated {n_components} components, Showing
675           PC1 vs PC2)",
676     labels={'x': f'PC1 ({pca.explained_variance_ratio_[0]:.2%})',
677             'y': f'PC2 ({pca.explained_variance_ratio_[1]:.2%})'}
678 )
679
680 pca_fig.update_layout(
681     title_font_family="serif",
682     title_font_color="blue",
683     font_family="serif",
684     font_size=14,
685     showlegend=True
686 )
687 pca_fig.update_xaxes(title_font_color="darkred", showgrid=True)
688 pca_fig.update_yaxes(title_font_color="darkred", showgrid=True)

```

```

689     return pca_fig
690
691
692 @app.callback(
693     Output("norm-result-text", "children"),
694     Input("clean-method", "value"),
695     Input("outlier-method", "value"),
696     Input("transform-method", "value"),
697     Input("range-slider", "value"),
698     Input("norm-feature", "value"),
699     Input("norm-method", "value")
700 )
701 def update_normality_test(clean, outlier, transform, range_val, feature,
702                           method):
703     if not feature:
704         return "Please select a feature."
705
706     df = get_processed_df(clean, outlier, transform, range_val)
707     data = df[feature].dropna()
708
709     if len(data) < 3:
710         return "Not enough data points to perform normality tests."
711
712     result_string = perform_normality_tests(data, feature, method)
713     return result_string
714
715
716 @app.callback(
717     Output("numeric-graph", "figure"),
718     Input("clean-method", "value"),
719     Input("outlier-method", "value"),
720     Input("transform-method", "value"),
721     Input("range-slider", "value"),
722     Input("num-f1", "value"),
723     Input("num-f2", "value"),
724     Input("num-f3", "value"),
725     Input("num-type", "value"),
726     Input("kde-palette", "value"),
727     Input("kde-width", "value"),
728     Input("graph-options", "value")
729 )
730 def update_numeric_plot(clean, outlier, transform, range_val, f1, f2, f3,
731                        plot_type, palette, width, options):

```

```

730 df = get_processed_df(clean, outlier, transform, range_val)
731 if df.empty: return make_warning_figure("No Data")
732
733 df_sub = df.sample(min(len(df), 5000), random_state=42)
734
735 if "Label" in df_sub.columns and hasattr(df_sub["Label"], 'cat'):
736     df_sub["Label"] = df_sub["Label"].astype(str)
737
738 if not f1: return make_warning_figure("Select Feature 1")
739
740 fig = go.Figure()
741
742 if plot_type == "line":
743     df_sorted = df_sub.sort_index()
744     fig = px.line(df_sorted, x=df_sorted.index, y=f1, title=f"Line Plot: {f1}")
745     fig.update_layout(xaxis_title="Index")
746     fig.update_traces(line=dict(width=width))
747
748 elif plot_type == "dist":
749     fig = px.histogram(df_sub, x=f1, title=f"Distribution Plot: {f1}")
750
751 elif plot_type == "pair":
752     cols = [c for c in [f1, f2, f3] if c]
753     if len(cols) < 2: return make_warning_figure("Select at least 2 features")
754     fig = px.scatter_matrix(df_sub, dimensions=cols, color="Label" if "Label" in df_sub.columns else None,
755                             title="Pair Plot")
756
757 elif plot_type == "heatmap":
758     cols = [c for c in [f1, f2, f3] if c]
759     if len(cols) < 2: return make_warning_figure("Select at least 2 features")
760     fig = px.density_heatmap(df_sub, x=f1, y=f2, title=f"Heatmap: {f1} vs {f2}")
761
762 elif plot_type == "hist_kde":
763     fig = px.histogram(df_sub, x=f1, marginal="violin", title=f"Histogram with KDE: {f1}")
764
765 elif plot_type == "qq":
766     data = df_sub[f1].dropna()

```

```

767     (osm, osr), _ = stats.probplot(data, dist="norm")
768     fig = go.Figure()
769     fig.add_trace(go.Scatter(x=osm, y=osr, mode='markers', name='Data'))
770     fig.add_trace(go.Scatter(x=[min(osm), max(osm)], y=[min(osm), max(osm)],
771                     mode='lines', line=dict(color='red'),
772                                     name='Normal Line'))
773     fig.update_layout(
774         title=f"QQ Plot: {f1}",
775         xaxis_title="Theoretical Quantiles",
776         yaxis_title="Sample Quantiles"
777     )
778
779 elif plot_type == "kde_custom":
780     data = df_sub[f1].dropna()
781     if len(data) > 1:
782         kde = stats.gaussian_kde(data)
783         x_range = np.linspace(data.min(), data.max(), 200)
784         y_val = kde(x_range)
785         color_map = {"blue": "#1f77b4", "red": "#d62728", "green": "#2ca02c", "purple": "#9467bd"}
786         line_color = color_map.get(palette, "blue")
787         fig = go.Figure()
788         fig.add_trace(go.Scatter(x=x_range, y=y_val, mode='lines', fill='tozero',
789                               line=dict(color=line_color, width=width),
790                               opacity=0.6, name="KDE"))
791         fig.update_layout(title=f"KDE Plot (Alpha=0.6, Width={width}): {f1}",
792                           xaxis_title=f"{f1}",
793                           yaxis_title="Density", )
794     else:
795         return make_warning_figure("Not enough data for KDE")
796
797 elif plot_type == "reg":
798     if not f2: return make_warning_figure("Select Feature 2")
799     fig = px.scatter(df_sub, x=f1, y=f2, trendline="ols", title=f"Regression: {f1} vs {f2}",
800                       trendline_color_override="red")
801
802 elif plot_type == "area":
803     if not f2: return make_warning_figure("Select Feature 2")
804     df_sorted = df_sub.sort_values(by=f1)
805     fig = px.area(df_sorted, x=f1, y=f2, title=f"Area Plot: {f1} vs {f2}")

```

```

804
805     elif plot_type == "joint":
806         if not f2: return make_warning_figure("Select Feature 2")
807         fig = px.scatter(df_sub, x=f1, y=f2, marginal_x="violin", marginal_y="violin",
808                           title="Joint Plot (Scatter + KDE-style Marginals)")
809
810     elif plot_type == "rug":
811         fig = px.strip(df_sub, x=f1, title=f"Rug Plot: {f1}")
812         fig.update_layout(
813             xaxis_title=f"{f1}",
814             yaxis_title="Occurrences",
815         )
816
817     elif plot_type == "3d":
818         if not f3: return make_warning_figure("Select Feature 3")
819         fig = px.scatter_3d(df_sub, x=f1, y=f2, z=f3, color="Label" if "Label"
820                             in df_sub.columns else None,
821                             title="3D Plot")
822
823     elif plot_type == "contour":
824         if not f2: return make_warning_figure("Select Feature 2")
825         fig = px.density_contour(df_sub, x=f1, y=f2, title=f"Contour Plot: {f1}
826             } vs {f2}")
827
828     elif plot_type == "cluster":
829         cols = [c for c in NUMERIC_COLS if c in df.columns][:10]
830         corr = df[cols].corr()
831         fig = px.imshow(corr, title="Cluster Map")
832
833     elif plot_type == "hexbin":
834         if not f2: return make_warning_figure("Select Feature 2")
835         fig = px.density_hexbin(df_sub, x=f1, y=f2, nbinsx=30, nbinsy=30,
836                               title=f"Hexbin Plot: {f1} vs {f2}")
837
838     else:
839         return make_warning_figure("Unknown Plot Type")
840
841     show_grid = 'grid' in options
842     show_legend = 'legend' in options
843     fig.update_layout(
844         title_font_family="serif",
845         title_font_color="blue",

```

```

843     font_family="serif",
844     showlegend=show_legend,
845     font_size=14,
846 )
847 fig.update_xaxes(title_font_color="darkred", showgrid=show_grid)
848 fig.update_yaxes(title_font_color="darkred", showgrid=show_grid)
849
850 return fig
851
852
853 @app.callback(
854     Output("cat-graph", "figure"),
855     Input("clean-method", "value"),
856     Input("outlier-method", "value"),
857     Input("transform-method", "value"),
858     Input("range-slider", "value"),
859     Input("cat-col", "value"),
860     Input("cat-num", "value"),
861     Input("cat-type", "value"),
862     Input("bar-mode", "value"),
863     Input("graph-options", "value")
864 )
865 def update_categorical_plot(clean, outlier, transform, range_val, cat, num,
866                             plot_type, bar_mode, options):
867     df = get_processed_df(clean, outlier, transform, range_val)
868     if df.empty: return make_warning_figure("No Data")
869
870     df_sub = df.sample(min(len(df), 5000), random_state=42)
871
872     if "Label" in df_sub.columns and hasattr(df_sub["Label"], 'cat'):
873         df_sub["Label"] = df_sub["Label"].astype(str)
874
875     if not cat: return make_warning_figure("Select Categorical Feature")
876
877     fig = go.Figure()
878
879     if plot_type == "bar":
880         if not num: return make_warning_figure("Select Numeric Feature")
881         group_cols = [cat]
882         if "Label" in df_sub.columns and cat != "Label":
883             group_cols.append("Label")
884         df_agg = df_sub.groupby(group_cols, observed=False)[num].mean().
885             reset_index()

```

```

884     color_col = "Label" if "Label" in df_sub.columns else None
885     fig = px.bar(df_agg, x=cat, y=num, color=color_col, barmode=bar_mode,
886                   title=f"Bar Plot ({bar_mode})")
887
888     elif plot_type == "count":
889         fig = px.histogram(df_sub, x=cat, color="Label" if "Label" in df_sub.
890                             columns else None, barmode=bar_mode,
891                             title=f"Count Plot: {cat}")
892
893     elif plot_type == "pie":
894         fig = px.pie(df_sub, names=cat, title=f"Pie Chart: {cat}")
895
896     elif plot_type == "box_multi":
897         if not num: return make_warning_figure("Select Numeric Feature")
898         fig = px.box(df_sub, x=cat, y=num, color="Label" if "Label" in df_sub.
899                     columns else None,
900                     title=f"Box Plot: {num} by {cat}")
901
902     elif plot_type == "boxen_multi":
903         if not num: return make_warning_figure("Select Numeric Feature")
904         fig = px.box(df_sub, x=cat, y=num, color="Label" if "Label" in df_sub.
905                     columns else None, points="all",
906                     title=f"Boxen Plot")
907
908     elif plot_type == "violin":
909         if not num: return make_warning_figure("Select Numeric Feature")
910         fig = px.violin(df_sub, x=cat, y=num, color="Label" if "Label" in
911                         df_sub.columns else None, box=True,
912                         title=f"Violin Plot")
913
914     elif plot_type == "strip":
915         if not num: return make_warning_figure("Select Numeric Feature")
916         fig = px.strip(df_sub, x=cat, y=num, color="Label" if "Label" in
917                         df_sub.columns else None, title=f"Strip Plot")
918
919     else:
920         return make_warning_figure("Unknown Plot Type")

```

```

920
921     show_grid = 'grid' in options
922     show_legend = 'legend' in options
923     fig.update_layout(
924         title_font_family="serif",
925         title_font_color="blue",
926         font_family="serif",
927         showlegend=show_legend,
928         font_size=14,
929     )
930     fig.update_xaxes(title_font_color="darkred", showgrid=show_grid)
931     fig.update_yaxes(title_font_color="darkred", showgrid=show_grid)
932
933     return fig
934
935
936 @app.callback(
937     Output("story-graph", "figure"),
938     Input("clean-method", "value"),
939     Input("outlier-method", "value"),
940     Input("transform-method", "value"),
941     Input("range-slider", "value")
942 )
943 def update_storytelling(clean, outlier, transform, range_val):
944     df = get_processed_df(clean, outlier, transform, range_val)
945
946     if df.empty:
947         return make_warning_figure("No data available with current settings.")
948
949     df_sub = df.sample(min(len(df), 3000), random_state=42)
950
951     if "Label" in df_sub.columns and hasattr(df_sub["Label"], 'cat'):
952         df_sub["Label"] = df_sub["Label"].astype(str)
953
954     fig = make_subplots(
955         rows=2, cols=2,
956         subplot_titles=("Attack Distribution", "Flow Duration by Label",
957                         "Traffic Volume (Bytes vs Packets)", "Average Packet
958                         Size"),
959         specs=[[{"type": "domain"}, {"type": "xy"}],
960                [{"type": "xy"}, {"type": "xy"}]]
961 )

```

```

962     if "Label" in df_sub.columns:
963         counts = df_sub["Label"].value_counts().reset_index()
964         counts.columns = ["Label", "count"]
965         fig.add_trace(go.Pie(labels=counts["Label"], values=counts["count"],
966                               name="Attacks"), row=1, col=1)
967
968         for label in df_sub["Label"].unique():
969             subset = df_sub[df_sub["Label"] == label]
970             fig.add_trace(go.Box(y=subset["Flow Duration"], name=str(label),
971                                  showlegend=False), row=1, col=2)
972
973             avg_size = df_sub.groupby("Label", observed=False)[["Average Packet
974             Size"]].mean().reset_index()
975             fig.add_trace(go.Bar(x=avg_size["Label"], y=avg_size["Average Packet
976             Size"], name="Avg Size", showlegend=False),
977                           row=2, col=2)
978
979         else:
980             fig.add_annotation(text="No Label Column", row=1, col=1)
981
982         fig.add_trace(go.Scatter(x=df_sub["Flow Packets/s"], y=df_sub["Flow Bytes/
983             s"],
984             mode='markers', marker=dict(size=5, opacity=0.5),
985             name="Traffic", showlegend=False), row=2, col=1)
986
987         fig.update_layout(height=800, title_text="Comprehensive Network Traffic
988             Analysis")
989         return fig
990
991
992
993
994
995 if __name__ == "__main__":
996     app.run(debug=False, host="0.0.0.0", port=8080)

```

17 References

1. Canadian Institute for Cybersecurity. (2017). *Intrusion detection evaluation dataset (CIC-IDS2017)*. University of New Brunswick. <https://www.unb.ca/cic/datasets/ids-2017.html>
2. Plotly. (2025). *Plotly Python Graphing Library*. <https://plotly.com/python/>
3. Plotly. (2025). *Dash Documentation & User Guide*. <https://dash.plotly.com/>
4. Sharafaldin, I., Lashkari, A. H., & Ghorbani, A. A. (2018). Toward Generating a New Intrusion Detection Dataset and Intrusion Traffic Characterization. *Proceedings of the 4th International Conference on Information Systems Security and Privacy (ICISSP)*.